

02616

Large-scale Modelling

Scaling and data-types

mpi4py

Scaling – how to

Scaling is an intrinsic component of understanding large scale applications

An application will have to explain it's limitations by means of scaling plots!

<https://hpc-wiki.info/hpc/Scaling>



Read every week!
Details here should
be 2nd nature!

Scaling: strong vs. weak

- ❑ How does the execution time go down for a fixed problem size by increasing the number of PUs?
 - ❑ Amdahl's law \Rightarrow speed-up, i.e. reduce time
 - ❑ also known as “strong scaling”
- ❑ How much can we increase the problem size by adding more PUs, keeping the execution time approx. constant?
 - ❑ Gustafson's law \Rightarrow scale-up, i.e. increase work
 - ❑ also known as “weak scaling”

Strong scaling

Fixed problem size, increase processors

1. Decide on problem size N
2. Get execution time (T) for 1, 2, 4, 8, ..., NP processors, same N
3. Plot NP vs. T (1) / T (NP)

A program has a serial fraction s and a parallel fraction p .

In the ideal world one would assume that:

- s is constant, irrespective of NP
- p is constant, and that the time of the parallel fraction becomes p/NP

So the execution time is

$$T(NP) \propto s + p/NP$$

Amdahl's law

scaling



$$S = \frac{T(1)}{T(NP)} \propto \frac{s + p}{s + p/NP} = \frac{1}{s + p/NP}$$

Weak scaling

Scale the problem and the number of processors by the same factor

1. Decide on initial problem size N
2. Get execution time (T) for 1, 2, 4, ..., NP processors, scale N with NP
(i.e. problem size gets doubled when doubling the number of processors)
3. Plot NP vs. $T(1) / T(NP)$



A program has a serial fraction s and a parallel fraction p .

In the ideal world one have:

- s is constant, irrespective of N
- p is constant, and the total parallel time in the scaled problem is p

So the scaled speedup is

$$S(NP) = \frac{\tau(NP)}{\tau(1)} = \frac{s + p \cdot NP}{s + p} = s + p \cdot NP$$

Efficiency




$$E(NP) = \frac{T(1)}{T(NP)}$$


Gustafson's law

Amdahl's vs Gustafson's law

Amdahl: fixed work



Gustafson: fixed work/PU



Amdahl's vs Gustafson's Law

❑ Amdahl's law

- ❑ Theoretical performance of an application with a ***fixed amount of parallel work*** given a particular number of Processing Units (PUs)

❑ Gustafson's Law:

- ❑ Theoretical performance of an application with a ***fixed amount of parallel work per PU*** given a particular number of PUs

Custom data-types in MPI

... on to something completely different ...

Custom data-types – contiguity

Consecutive memory

0	1	2	3	4	5	6	7	8	9
10	11	12	13	14	15	16	17	18	19


`MPI.Send([buf, N, dtype], ...)`

Sends N consecutive elements, each of a byte size determined by `dtype`. ($N \cdot \text{dtype.itemsize}$)

Recall week 4 where we sent a complex array with a float data type, resulting in only sending half of the complex number!

Boiling it down: a data-type tells MPI several details:

- An initial byte offset from the memory segment a user passes
- How many bytes a single element really is (float64 == 8 bytes)
- Where an element starts
- Where an element stops



Custom data-types – contiguity

Recall – sender and receiver need not use same data-type

Sending rank :
2 of (2*double, 2*hole)

0	1	2	3	4	5	6	7	8	9
10	11	12	13	14	15	16	17	18	19

Receiving rank:
4 of (1*double) or
2 of (2*double) or
1 of (4*double)

0	1	2	3	4	5	6	7	8	9
10	11	12	13	14	15	16	17	18	19

It is just a byte-stream!

Custom data-types – contiguity

Data-type
Datatype.Create_contiguous(int count)

When you have created your data-type; tell MPI to *use it!*


Datatype.Commit() # enables it to be used by MPI
Datatype.Free() # remove it from MPI

Easy to forget!

Simplest when packing data.

Used: Always sending N of chunks of specific size

E.g. a coordinate system (xyz).



```
dt = MPI.DOUBLE.Create_contiguous(3)  
dt.Commit() # necessary!  
Comm.Send([xyz, 2, dt])
```


Custom data-types – contiguity

Data-type
Datatype.Create_indexed(blocklengths, displacements)
Number of contiguous elements in each block
Displacement of DataType's size for each block


Datatype.Commit() # enables it to be used by MPI
Datatype.Free() # remove it from MPI

Easy to forget!

```
dt1 = MPI.DOUBLE.Create_indexed([2, 2], [0, 4])
```



```
dt2 = dt1.Create_indexed([1, 1], [0, 2])
```



Custom data-types – contiguity

Number of contiguous elements in each block
MPI.Datatype.Create_struct(blocklengths, displacements, dtypes)
Dtypes defining extent through blocklengths

```
dt1 = MPI.Datatype.Create_struct([2, 2], [0, 20], [MPI.DOUBLE, MPI.INT])
```

- use 2 doubles (16 bytes)
- jump to 20 byte
- use 2 integers (8 bytes)

An element is 4 bytes

0	1	2	3	4	5	6	7	8	9
10	11	12	13	14	15	16	17	18	19

Custom data-types – holes


Lower bound (in BYTES)

MPI.Datatype.Create_resized(lb, ub)

Upper bound of the data-types extent

```
dt = MPI.INT.Create_contiguous(2).Create_resized(0, 24)
```

- use 2 integers (8 bytes)
- the data-type stops at 24 byte, next element starts at 25 bytes



Send/Recv custom data-types

They behave as the normal dtypes, but mpi4py does not auto calculate the number of dtypes it will send, so always use the extended bufferspec.

mpi4py.typing.BufSpec

Buffer specification.

- `Buffer`
- `Tuple[Buffer , Count]`
- `Tuple[Buffer , TypeSpec]`
- `Tuple[Buffer , Count , TypeSpec]`
- `Tuple[Bottom , Count , Datatype]`

```
from mpi4py import MPI
C = MPI.COMM_WORLD
xyz = np.empty([N, 3])
dt = MPI.DOUBLE.Create_contiguous(3)
```

- `C.Send(xyz, 1)`
- `C.Send((a, N, dt), 1)`

Recall `Get_count`, can also be used with custom data-types!

```
C.Recv(buf, ..., status=status)
```

```
assert status.Get_count(dt) == N
```

```
assert status.Get_elements(dt) == N * 3
```

Count of dtype

Count of *basic* elements in dtype

Custom data-types

It is easy to create complicated data-types, customization is huge!

- Use them as much as you can, they will generally save you a temporary buffer when you do simple send/recv.
- They are local, so there is very little overhead
- Use `status.Get_count` to get count of data-types
- Use them for *boundary conditions* where a “halo” region needs to be communicated.
- They can also describe *sub-arrays* (see `MPI_Type_create_subarray`) (not covered! However, can be really useful when doing parallel IO in next weeks topic)
- They can reduce some complexity when you are sending constant sub-divisions i.e. coordinates ($N * 3 \Rightarrow$ dtype of length 3, send N)