

02616

Large-scale Modelling

Recap from Week 01

- ❑ parallel architectures:
 - ❑ UMA, NUMA, cc-NUMA
- ❑ first steps in MPI:
 - ❑ “Hello World” (no communication)
 - ❑ “Hello Neighbour” (communication)
 - ❑ Ring (communication)
- ❑ everything on a single node

Topics for Week 02

- ❑ Memory bandwidth and performance
 - ❑ “NUMA in action”
- ❑ Networks
 - ❑ cluster setup
 - ❑ Infiniband and Ethernet
 - ❑ bandwidth and latency
- ❑ MPI jobs on multiple nodes
 - ❑ multi-node batch jobs
 - ❑ process placement and binding

Memory bandwidth

Single-core peak performance:

- ❑ The TPP is easy to calculate: let us assume
 - ❑ 2 GHz CPU/core
 - ❑ 4 Flops per clock cycle
 - ❑ \Rightarrow TPP: 8 GFlop/s (per core!)
- ❑ To obtain that peak performance, we need to be able to feed the core with the right amount of data!

Memory bandwidth

How much data is that?

- ❑ 4 floating point operations (add, mult) need 8 floating numbers \Rightarrow 64 bytes (double precision)
- ❑ That is 64 bytes / clock cycle or 128 GB/s
 - ❑ 128 GB is the content of more than 27 DVDs!!!
- ❑ But what is the memory bandwidth in modern machines?

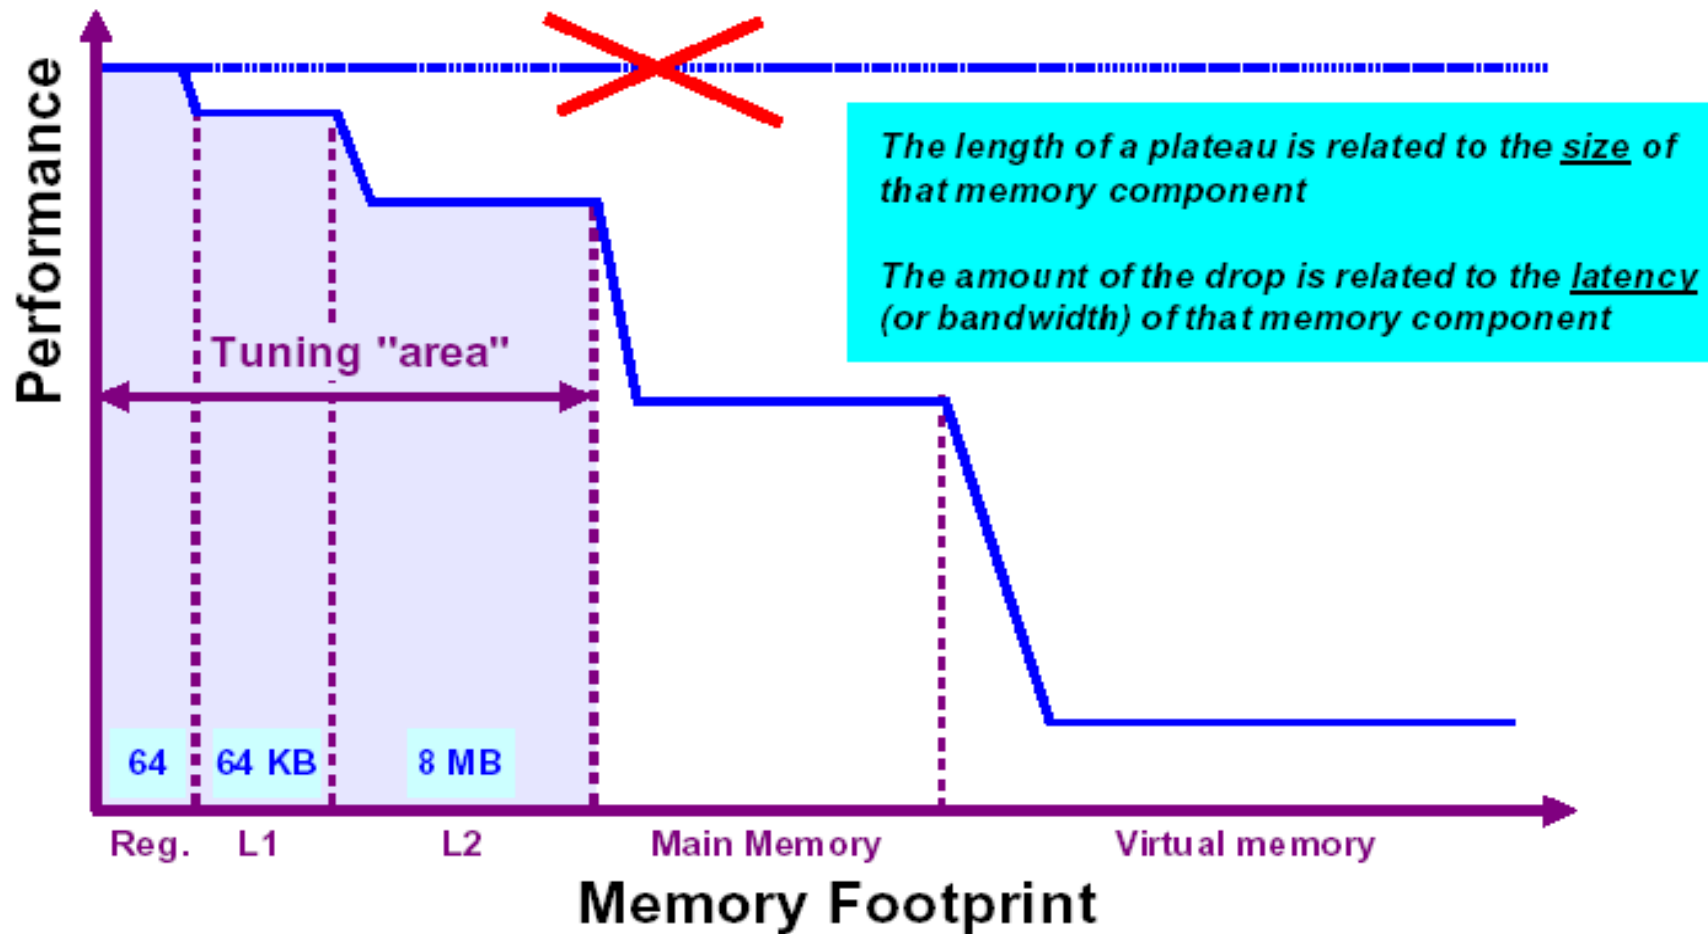
Memory bandwidth

Memory bandwidth of the Xeon E5-2650 v4 family

- ❑ equipped with DDR4-2400 DIMMS:
 - ❑ 2400 MT/s \Rightarrow 19.2 GB/s per memory channel
 - ❑ 76.8 GB/s maximum (all 4 channels equipped)
- ❑ with slower DDR4-1600 DIMMS:
 - ❑ max. 51.2 GB/s
- ❑ that is the bandwidth per CPU socket
 - ❑ ... but each CPU has 12 cores!
 - ❑ \Rightarrow the memory bandwidth per core is less!!!!

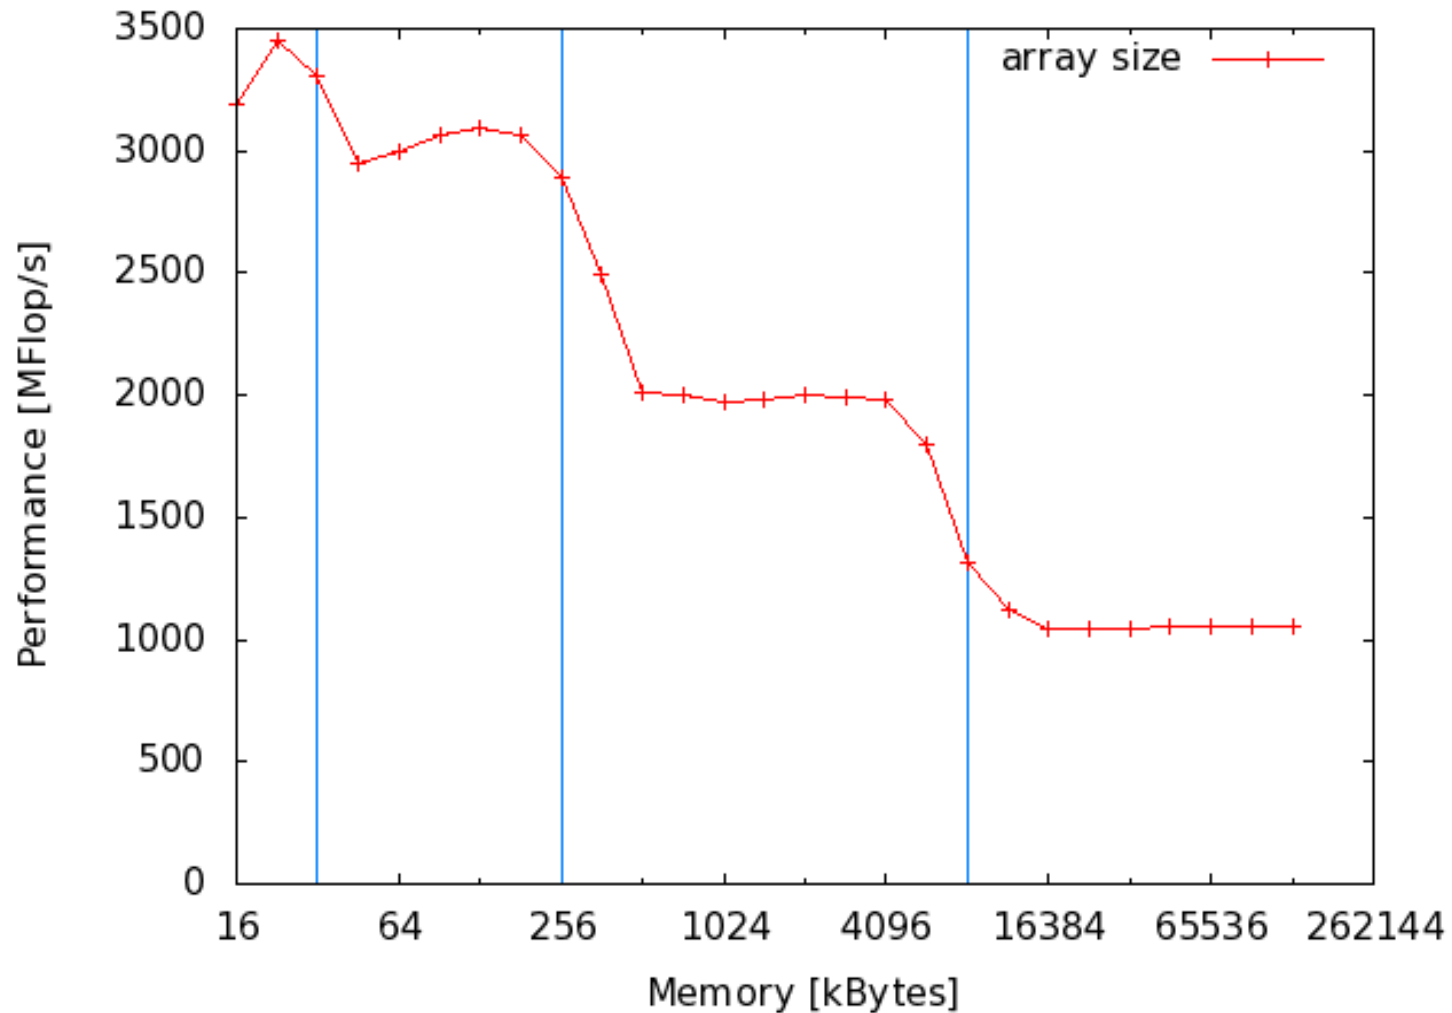
Memory bandwidth

Performance staircase (single core):



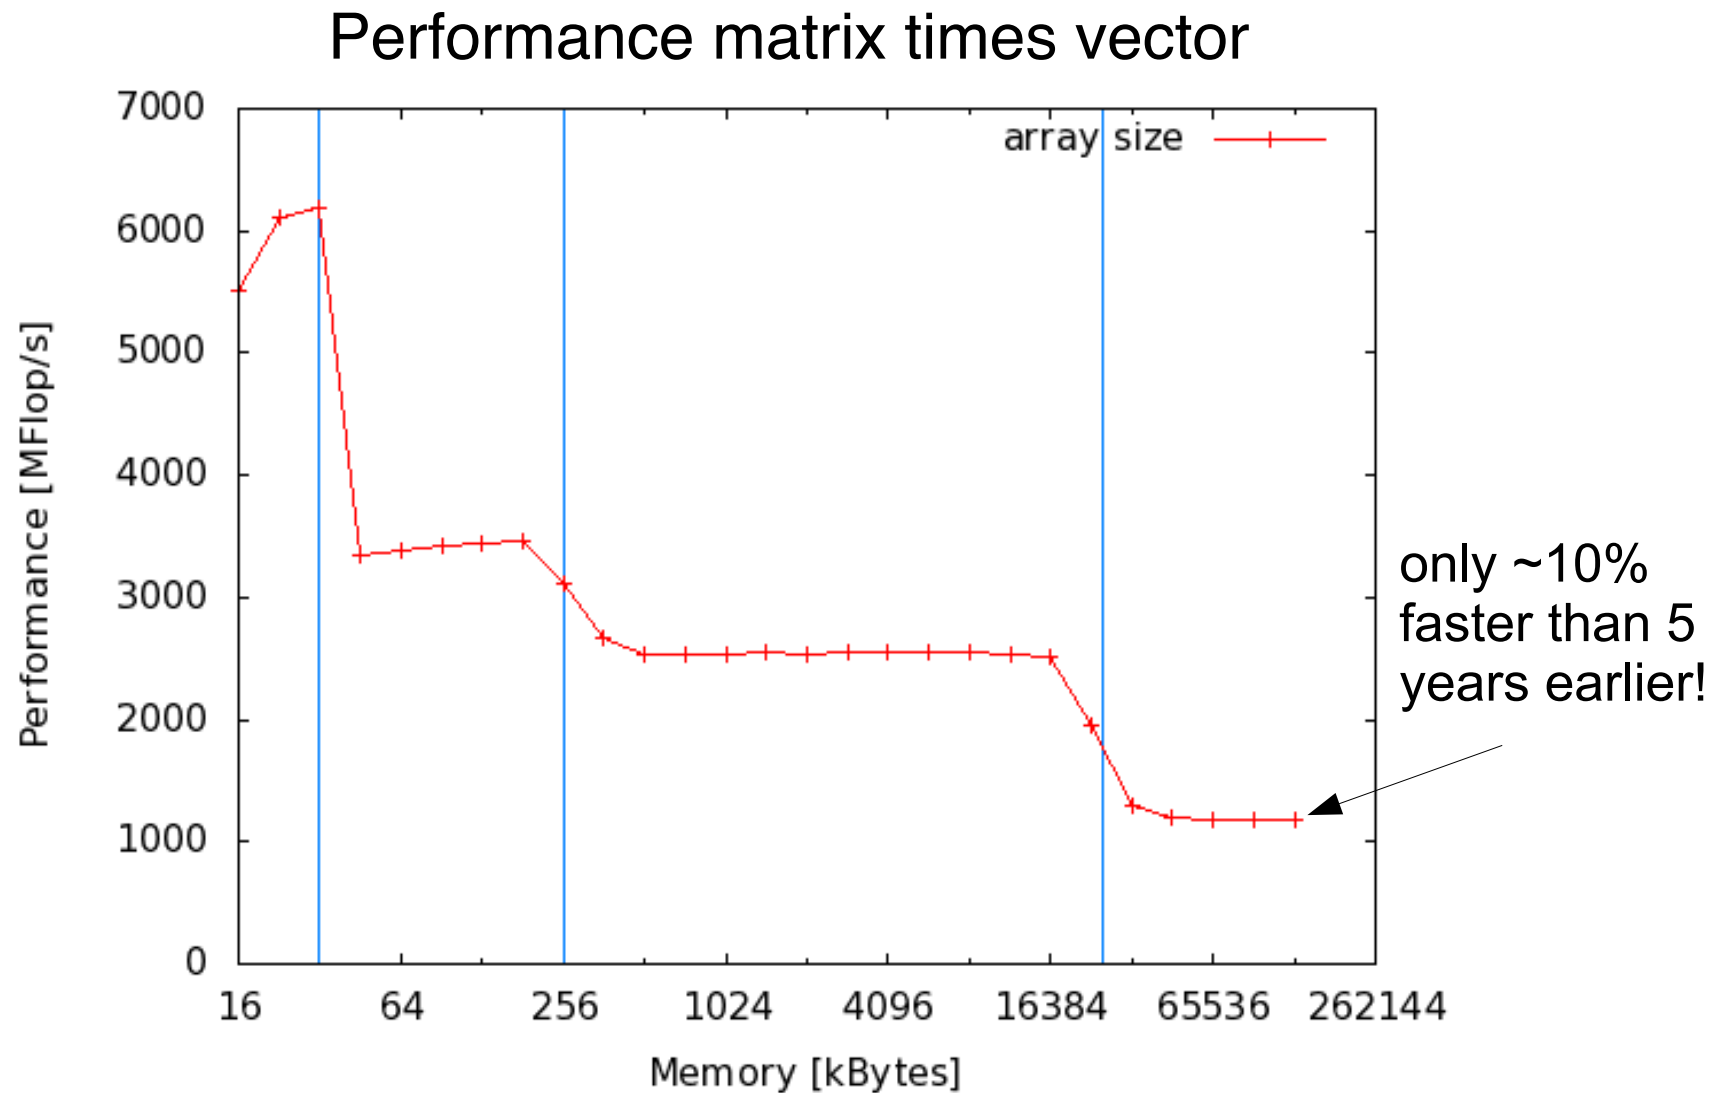
Memory bandwidth

Performance matrix times vector



Intel Xeon X5550, 2.67 GHz, 8 MB L3 cache (release: Q1/2009)

Memory bandwidth



Intel Xeon E5-2660v3, 2.66 GHz, 25 MB L3 (release: Q3/2014)

Memory bandwidth

Compute-bound vs memory-bound

❑ Compute-bound:

- ❑ the number of flops is higher than the number of mem-ops (load/store)
- ❑ example: matrix times matrix

❑ Memory bound:

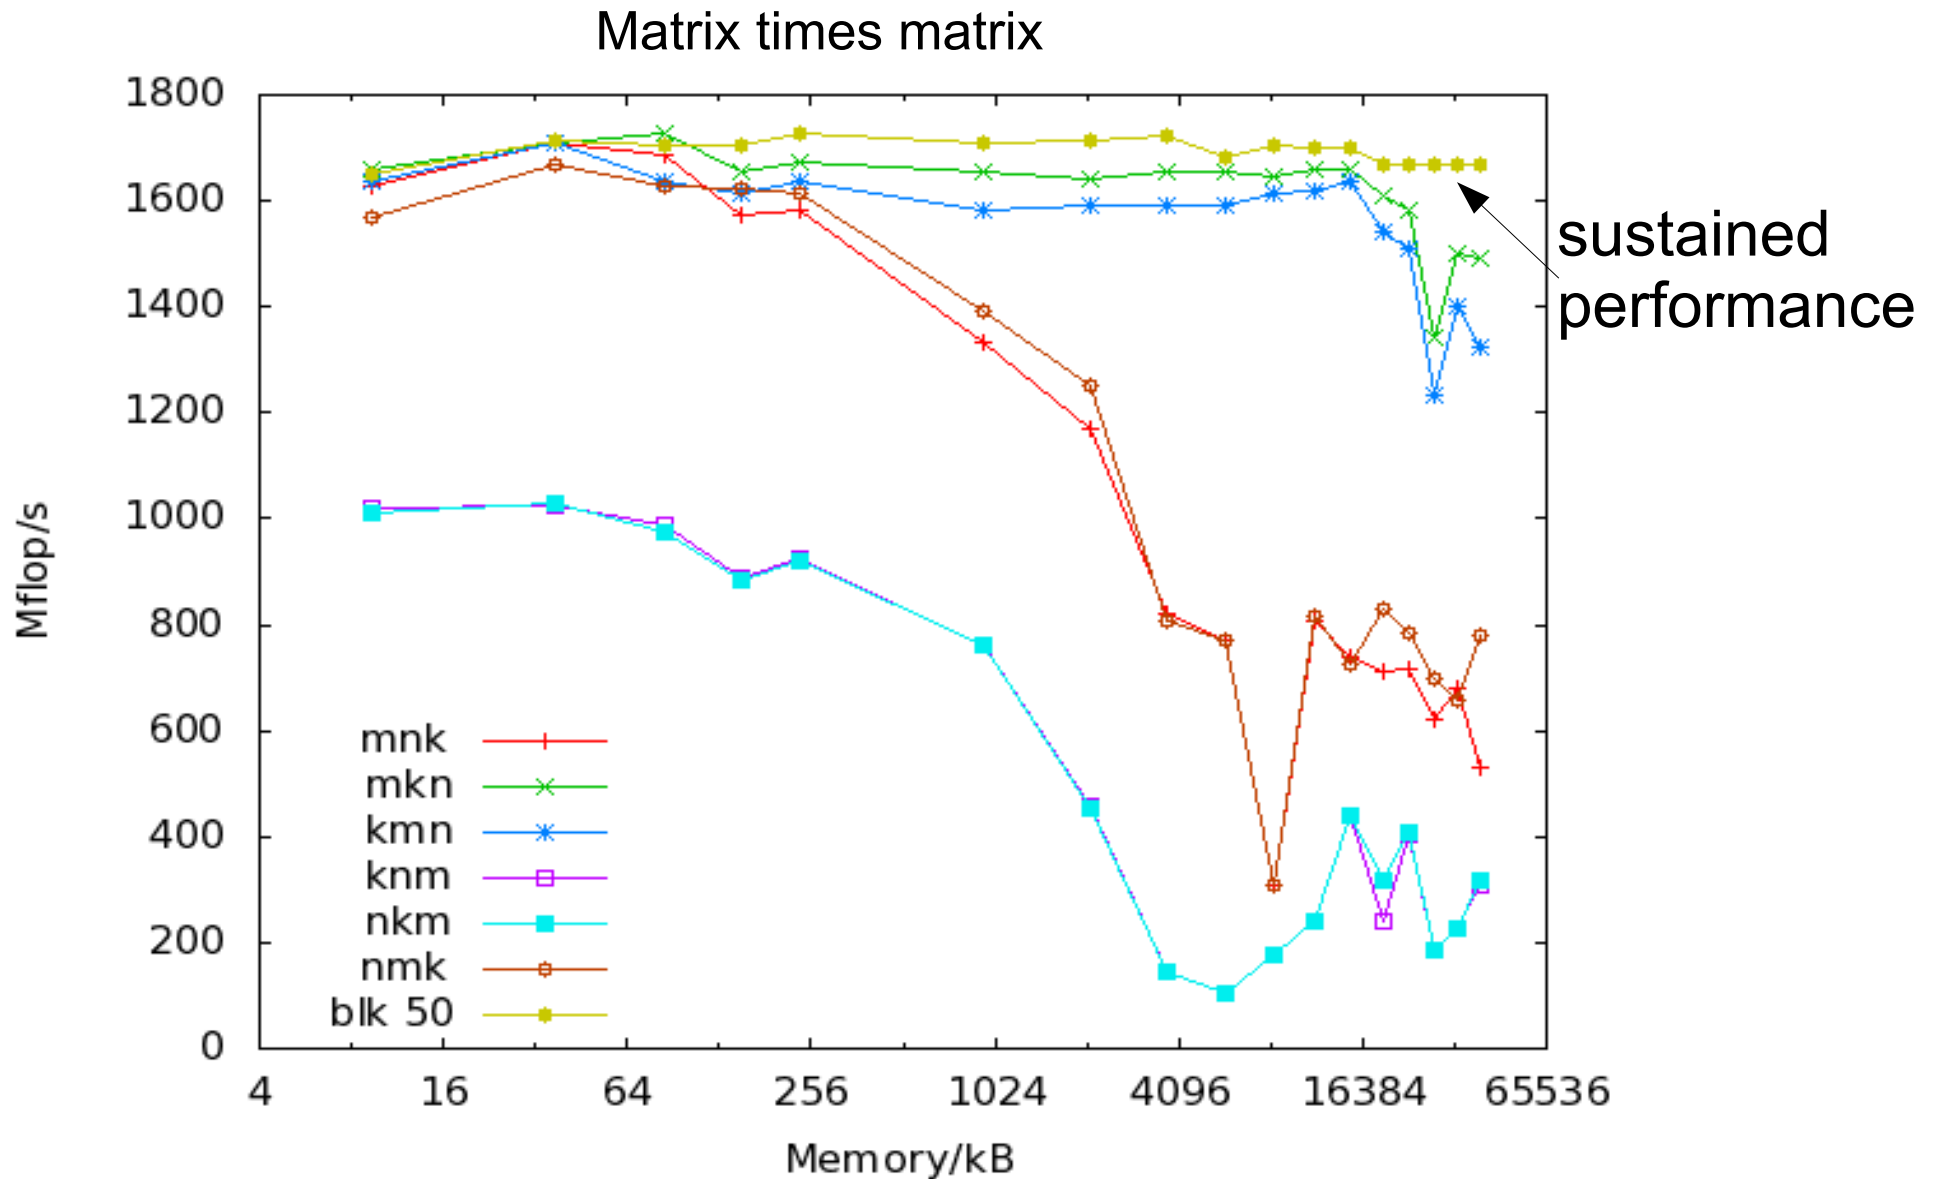
- ❑ the number of mem-ops is dominating the flops
- ❑ example: matrix times vector

Memory bandwidth

Compute-bound vs memory-bound

- ❑ but ...
- ❑ for large problems, even matrix times matrix gets dominated by the memory operations, it turns into a memory bound problem
- ❑ we know how to solve that – e.g. by blocking algorithms, to keep the problem compute bound

Memory bandwidth



Memory bandwidth

Scaling of memory-bound applications

- ❑ What happens with memory bound applications on multi-core systems?
- ❑ N cores – one CPU, e.g. $N = 4$
- ❑ N threads execute a memory-bound kernel
 - ❑ \Rightarrow all N threads fight for the same mem-bw
 - ❑ \Rightarrow this is a performance bottleneck
 - ❑ \Rightarrow this has a negative effect on scaling (speed-up)

Memory bandwidth

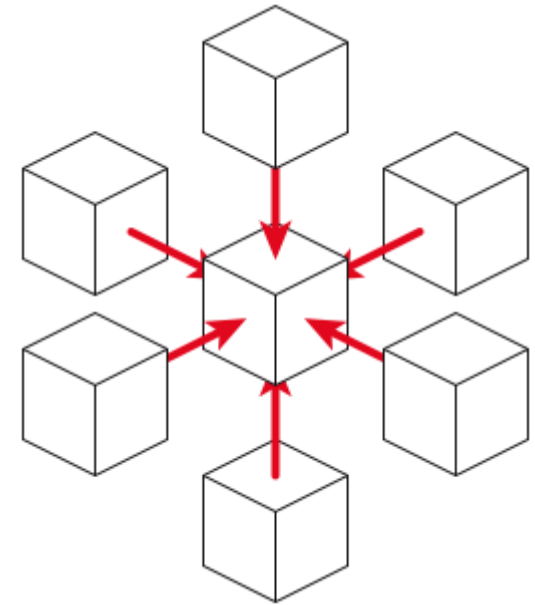
Scaling of memory-bound applications (cont'd)

- ❑ this effect is visible for all parallel applications, regardless of the programming model, i.e.
 - ❑ distributed applications (e.g. MPI)
 - ❑ shared-memory applications (e.g. OpenMP)
- ❑ adding more cores to the CPU makes the situation even worse!
- ❑ adding more sockets helps:
 - ❑ increased memory bandwidth
 - ❑ the overall performance is bounded by the max. memory bandwidth (saturation)

Memory bandwidth

Example (from the 02614 course):

- ❑ 3D Poisson problem: N^3 grid points
- ❑ 7 point stencil
- ❑ 8 mem-ops per stencil update
- ❑ memory bound problem
- ❑ data: double precision (64 bit)
- ❑ $N = 1000$
- ❑ memory usage: 22.5 GB
- ❑ OpenMP parallel implementation



Memory bandwidth

Example (from the 02614 course):

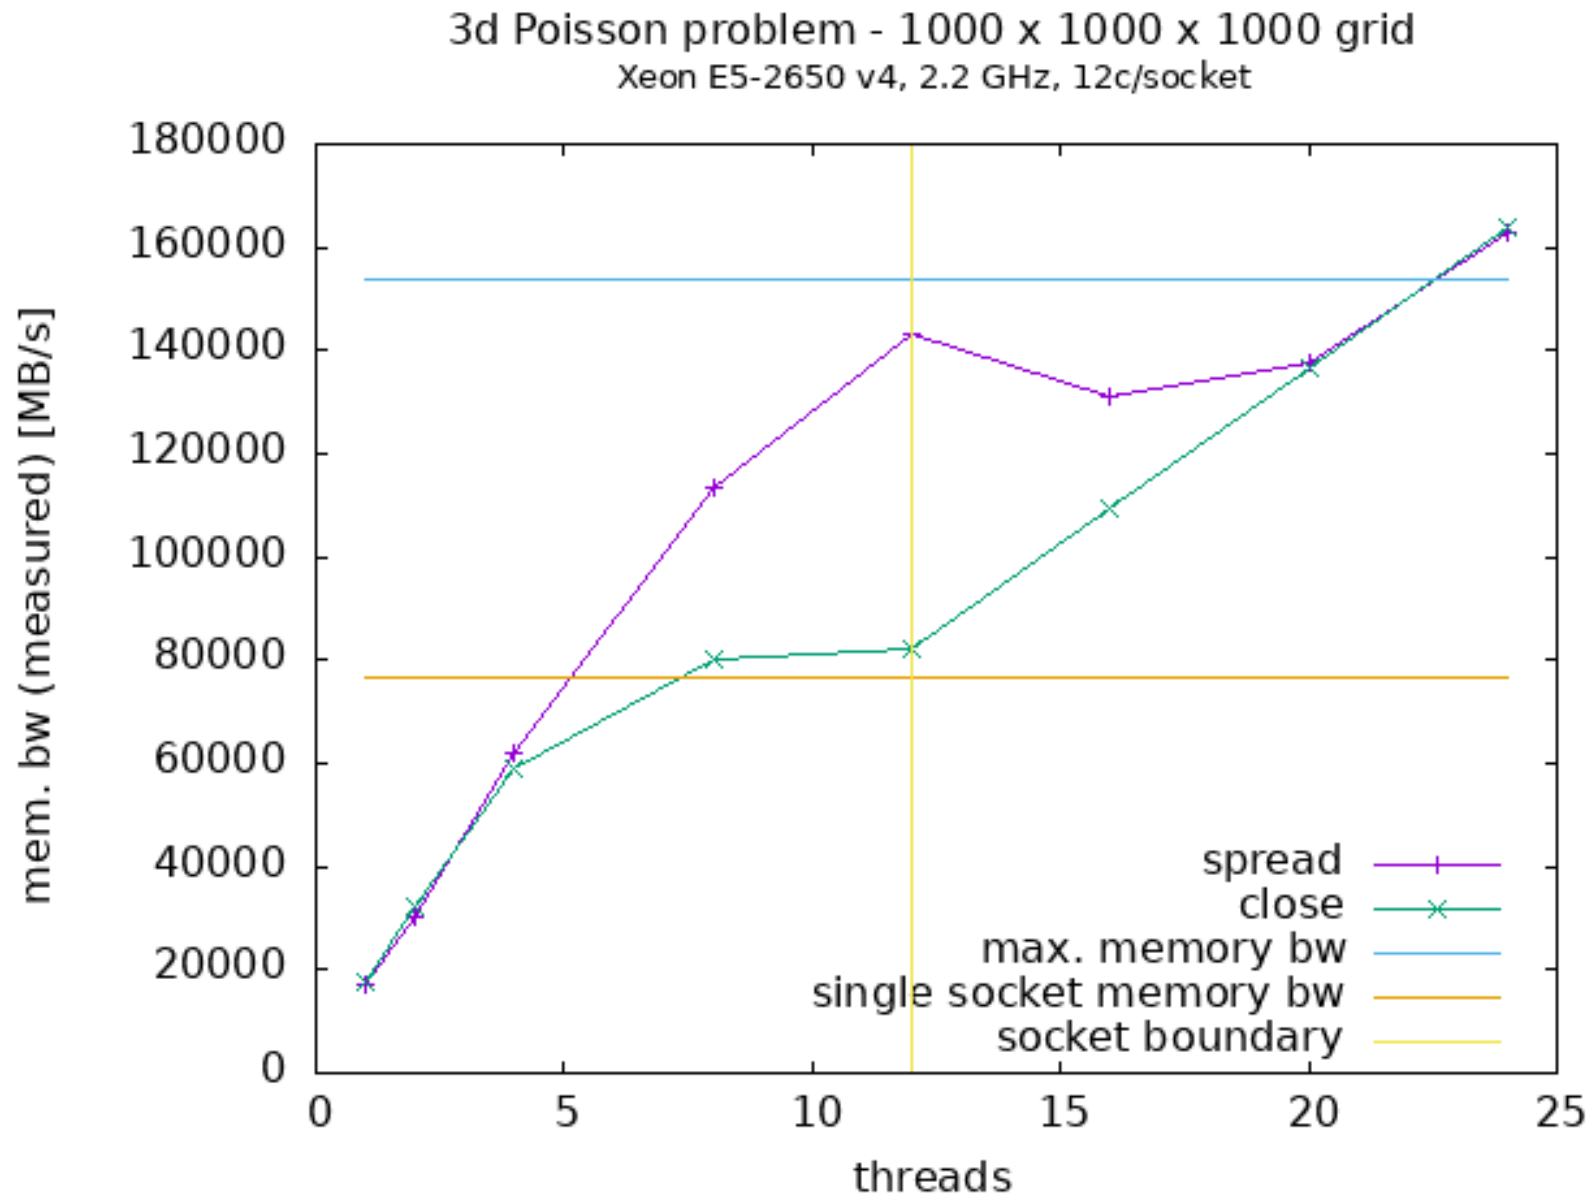
- ❑ 3D Poisson problem: 1000^3 grid points
- ❑ Hardware:
 - ❑ 2x Intel Xeon E5-2650 v4, 2.2 GHz
 - ❑ L3 size: 30MB
 - ❑ 256 GB RAM (128 GB per socket)
 - ❑ single core TPP: 35.2 Gflop/s
 - ❑ 12 cores per socket
 - ❑ memory bandwidth per socket: 76.8 GB/s

Memory bandwidth

How did we do?

- ❑ we can measure the number of stencil updates (lup) per second: Mlup/s
- ❑ each lup involves 8 memory operations:
 - ❑ 7 loads and 1 store
 - ❑ uncertainty: loads and stores can overlap
- ❑ our data type is double: 8 bytes
 - ❑ \Rightarrow 64 bytes per lup
- ❑ measured memory bandwidth:
 - ❑ $\text{MB/s} = \text{Mlup/s} * 64 \text{ bytes}$

Memory bandwidth



Memory bandwidth

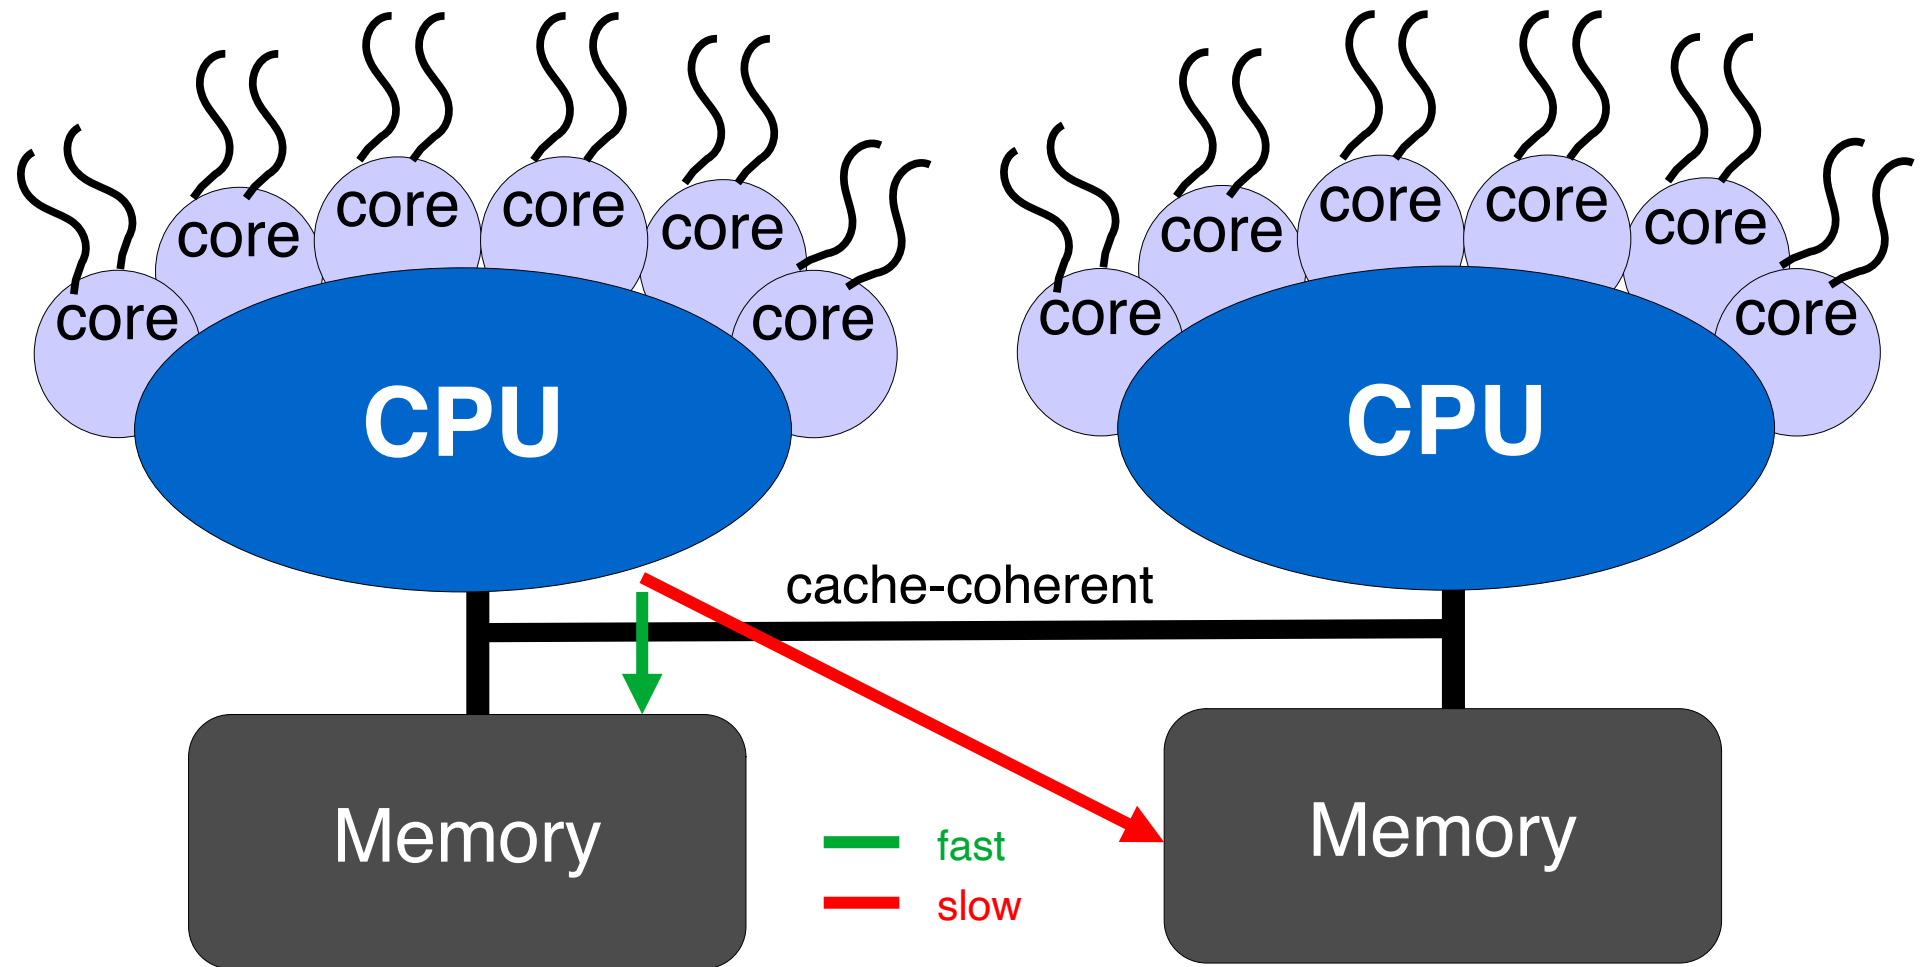
❑ Lessons learned:

- ❑ using a single socket (close distribution), we are already limited by the socket memory bandwidth when using 8 threads on the 12 cores
- ❑ spreading the threads over both sockets is beneficial for up to 12 threads, i.e. 6 threads per socket
- ❑ for larger number of threads, e.g. 20 and 24 threads (the full node), the two distributions approach each other.
- ❑ similar behaviour can be observed when using an MPI implementation of the problem.

Networks

Cluster of multi-socket machines

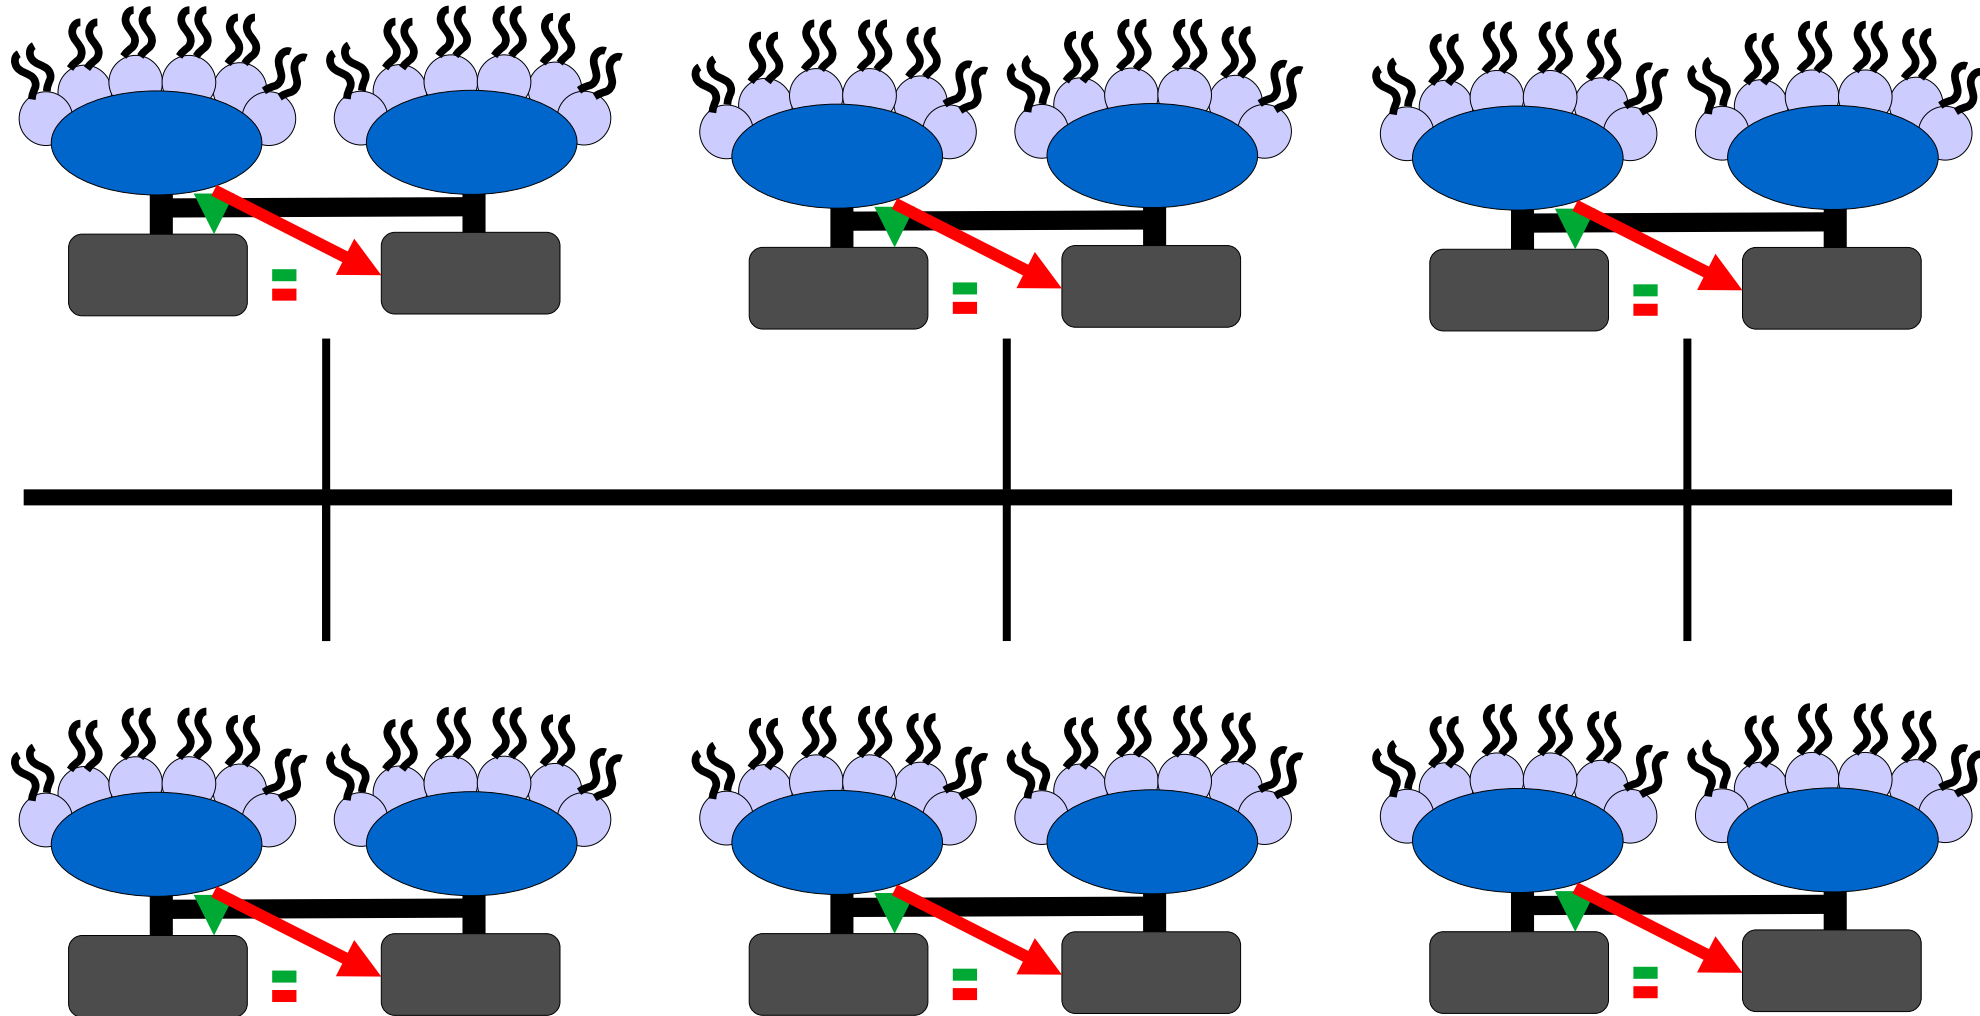
A typical multi-core server



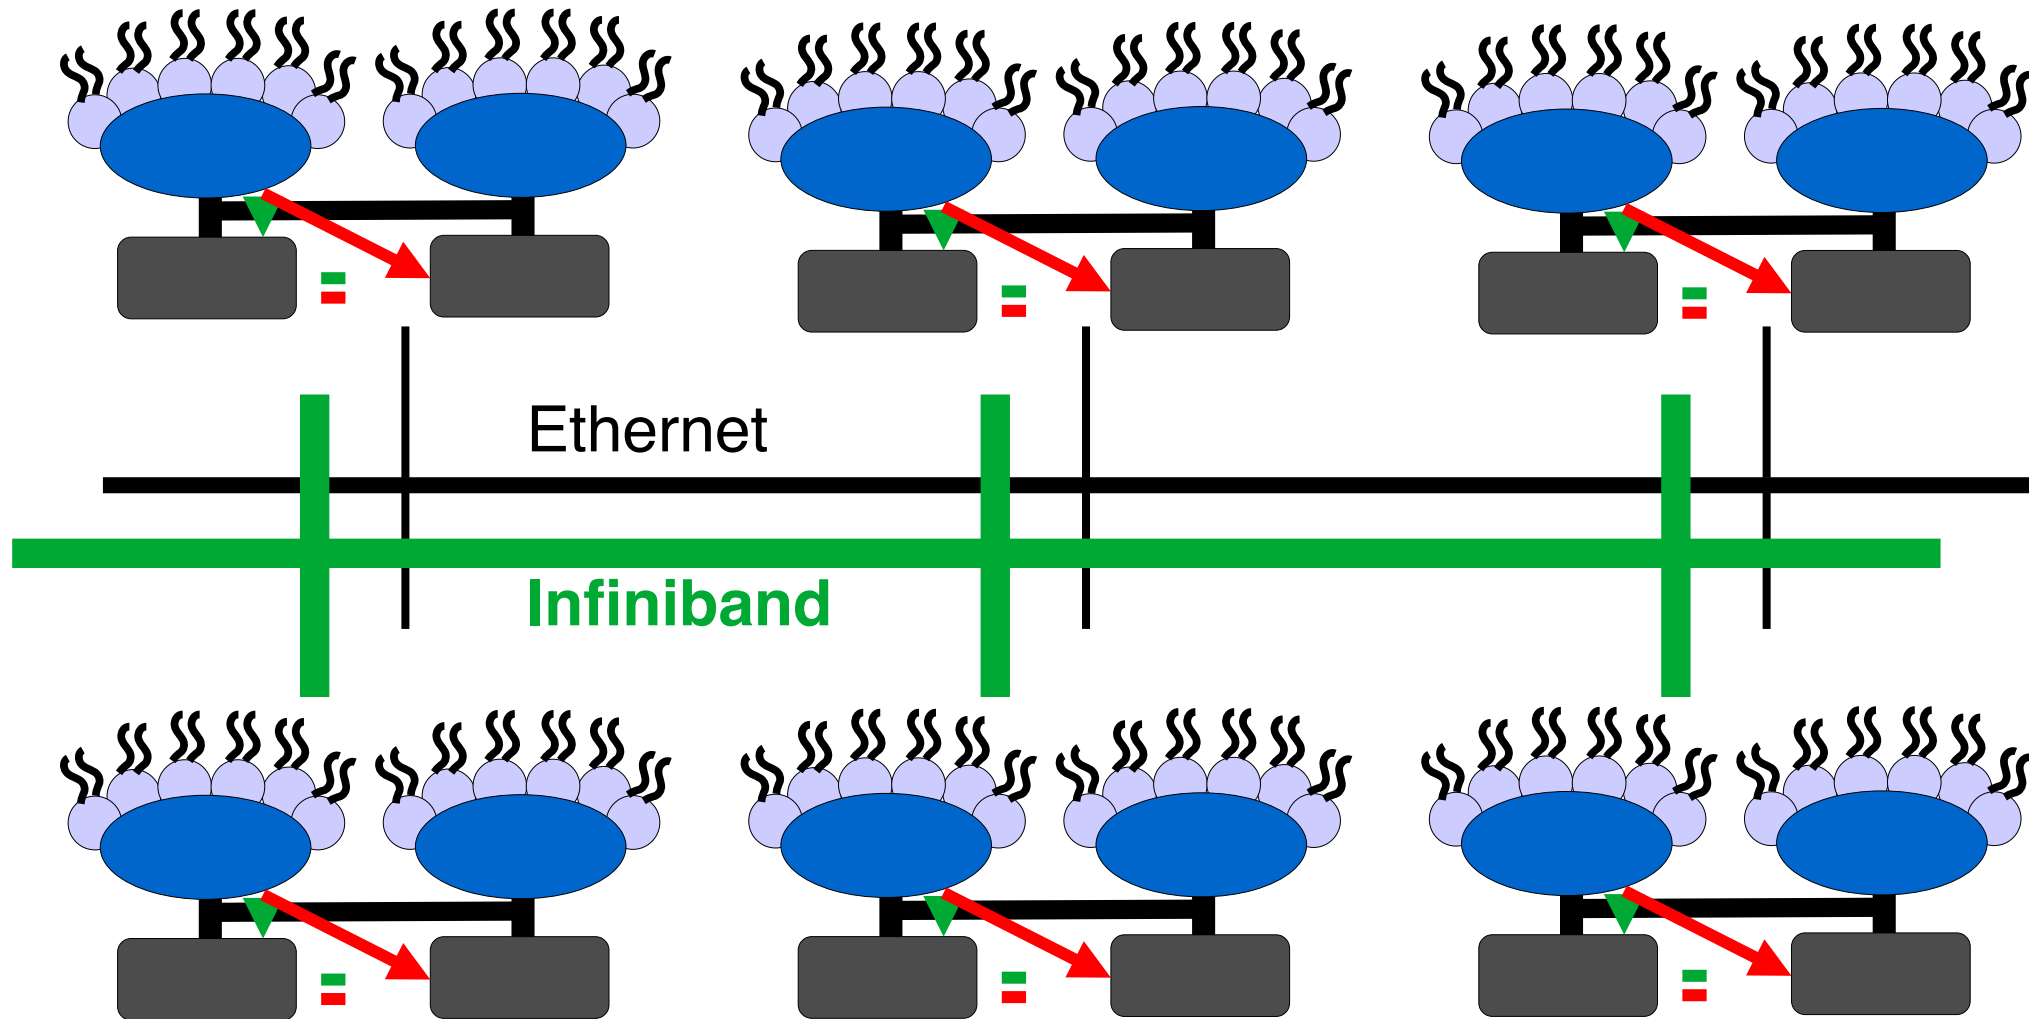
a 2-socket, 12-core, 24-(hyper-)threads server \Rightarrow 24 logical CPUs

Note: we do not use hyperthreading in our setup!

Cluster with a single network



Cluster with two networks



Why a fast network?

- ❑ for HPC applications, we want to have
 - ❑ a high network bandwidth
 - ❑ low latency
 - ❑ => minimizing the communication overhead
- ❑ this can be achieved by networks like Infiniband (IB), OmniPath (OPA), etc
- ❑ other advantages of IB
 - ❑ RDMA: remote direct memory access – doesn't involve the CPU when accessing memory on a remote node
 - ❑ GPU Direct: RDMA between a local and remote GPU

Why a fast network?

Infiniband specifications:

- ❑ there are several generations of IB
- ❑ the latest generations provide
 - ❑ up to 400 Gb/s
 - ❑ 3-5 microseconds latency
- ❑ in our setup:
 - ❑ FDR IB: up 56 Gb/s
 - ❑ compare to Ethernet: either 1 Gb/s or 10 Gb/s
 - ❑ today's lab: measure the difference in bandwidth and latency between IB and Ethernet

Fun fact

What is the memory bandwidth of our brain?

- ❑ It's really slow: 10 bit/s!
- ❑ Read more about it
 - ❑ "The unbearable slowness of being: Why do we live at 10 bits/s?"

The “real” world

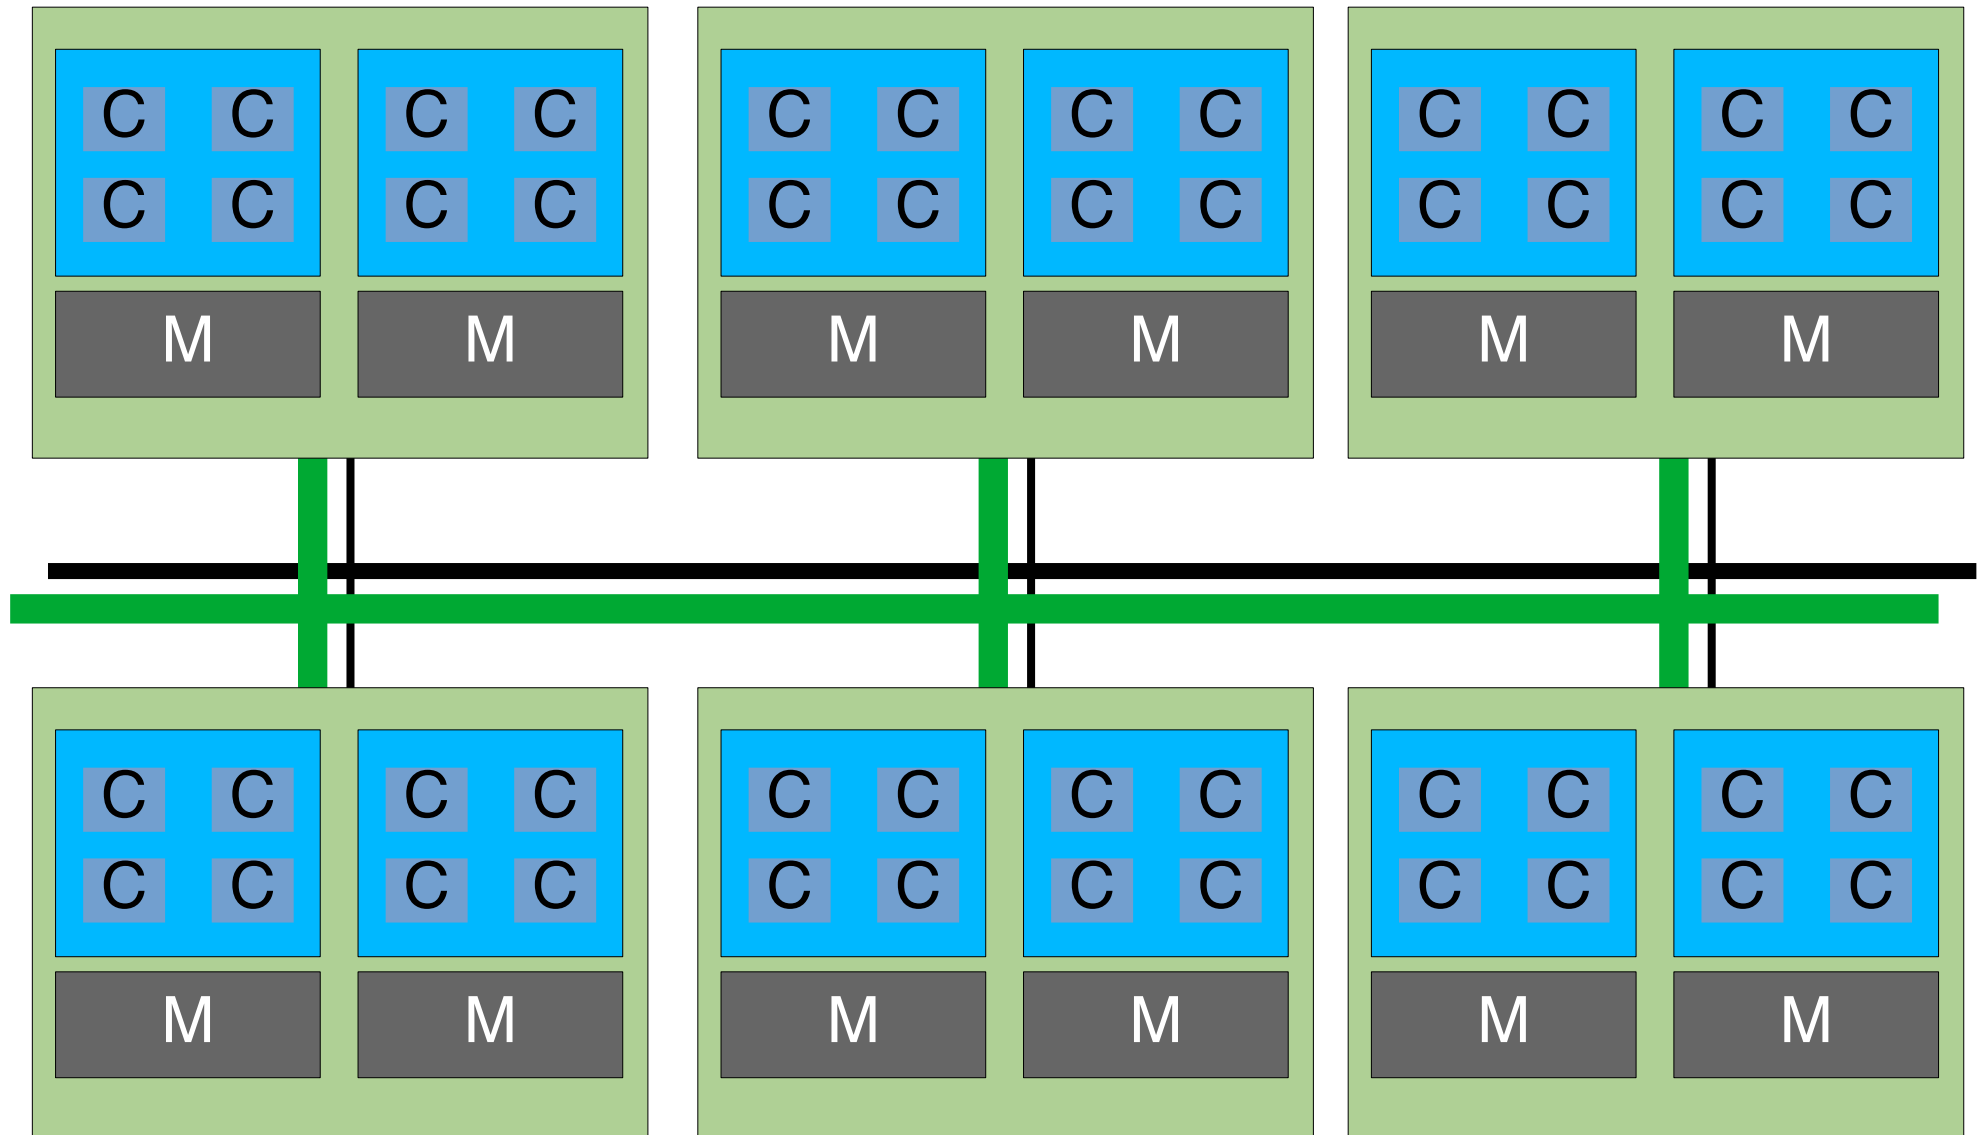
Putting it all together

The “real” world

What we have learned so far:

- ❑ The limiting factor in many applications is bandwidth:
 - ❑ memory bandwidth in a single node
 - ❑ network bandwidth on a network of nodes
- ❑ Now let us put those pieces together:
 - ❑ a cluster of multi-socket (cc-NUMA) systems,
 - ❑ connected by both Infiniband and Ethernet

The “real” world



The “real” world

- ❑ run an MPI application
 - ❑ with N ranks (single-threaded)
 - ❑ on M nodes (dual-socket)
 - ❑ each node has $2 \cdot P$ cores (P cores/socket)
- ❑ we will have $2 \cdot P \cdot M$ locations to place our N ranks on ($N \leq 2 \cdot P \cdot M$)
- ❑ How does the placement influence the performance of our application?
- ❑ How can we control it?
 - ❑ interplay between MPI and scheduler (LSF)

The “real” world

Specific example:

- ❑ $N = 8$ ranks
- ❑ $M = 4$ nodes
- ❑ $P = 4$ cores/socket (8 cores per node)
- ❑ 32 possible locations for the 8 ranks

The “real” world

Scenario 1:

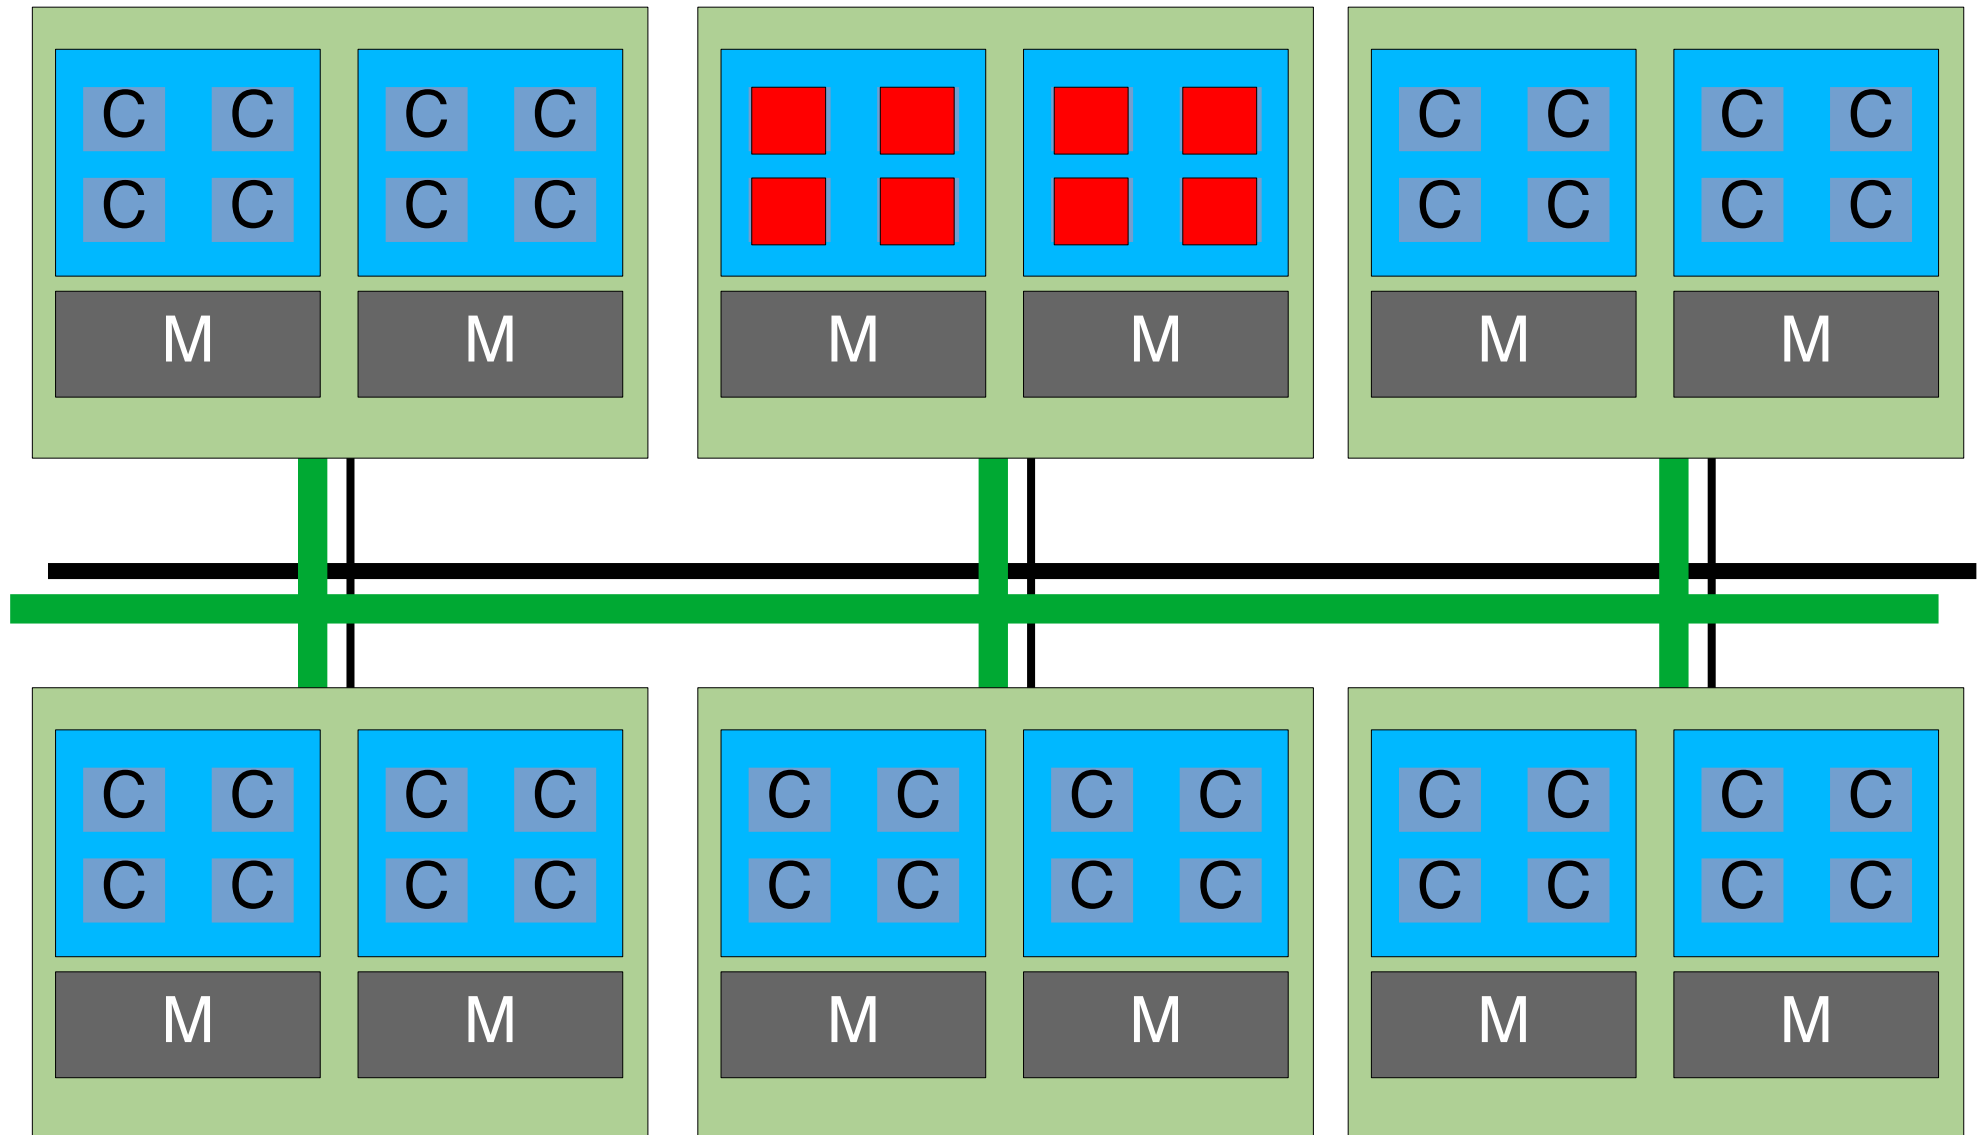
- ❑ all 8 ranks on a single node

```
#BSUB -J MPIapp
#BSUB -o MPIapp_%J.out
#BSUB -n 8
#BSUB -R "span[hosts=1]"

module load mpi/<version>

mpirun ./myapp
```


The “real” world



The “real” world

Scenario 2:

- ❑ use two nodes, 4 ranks per node

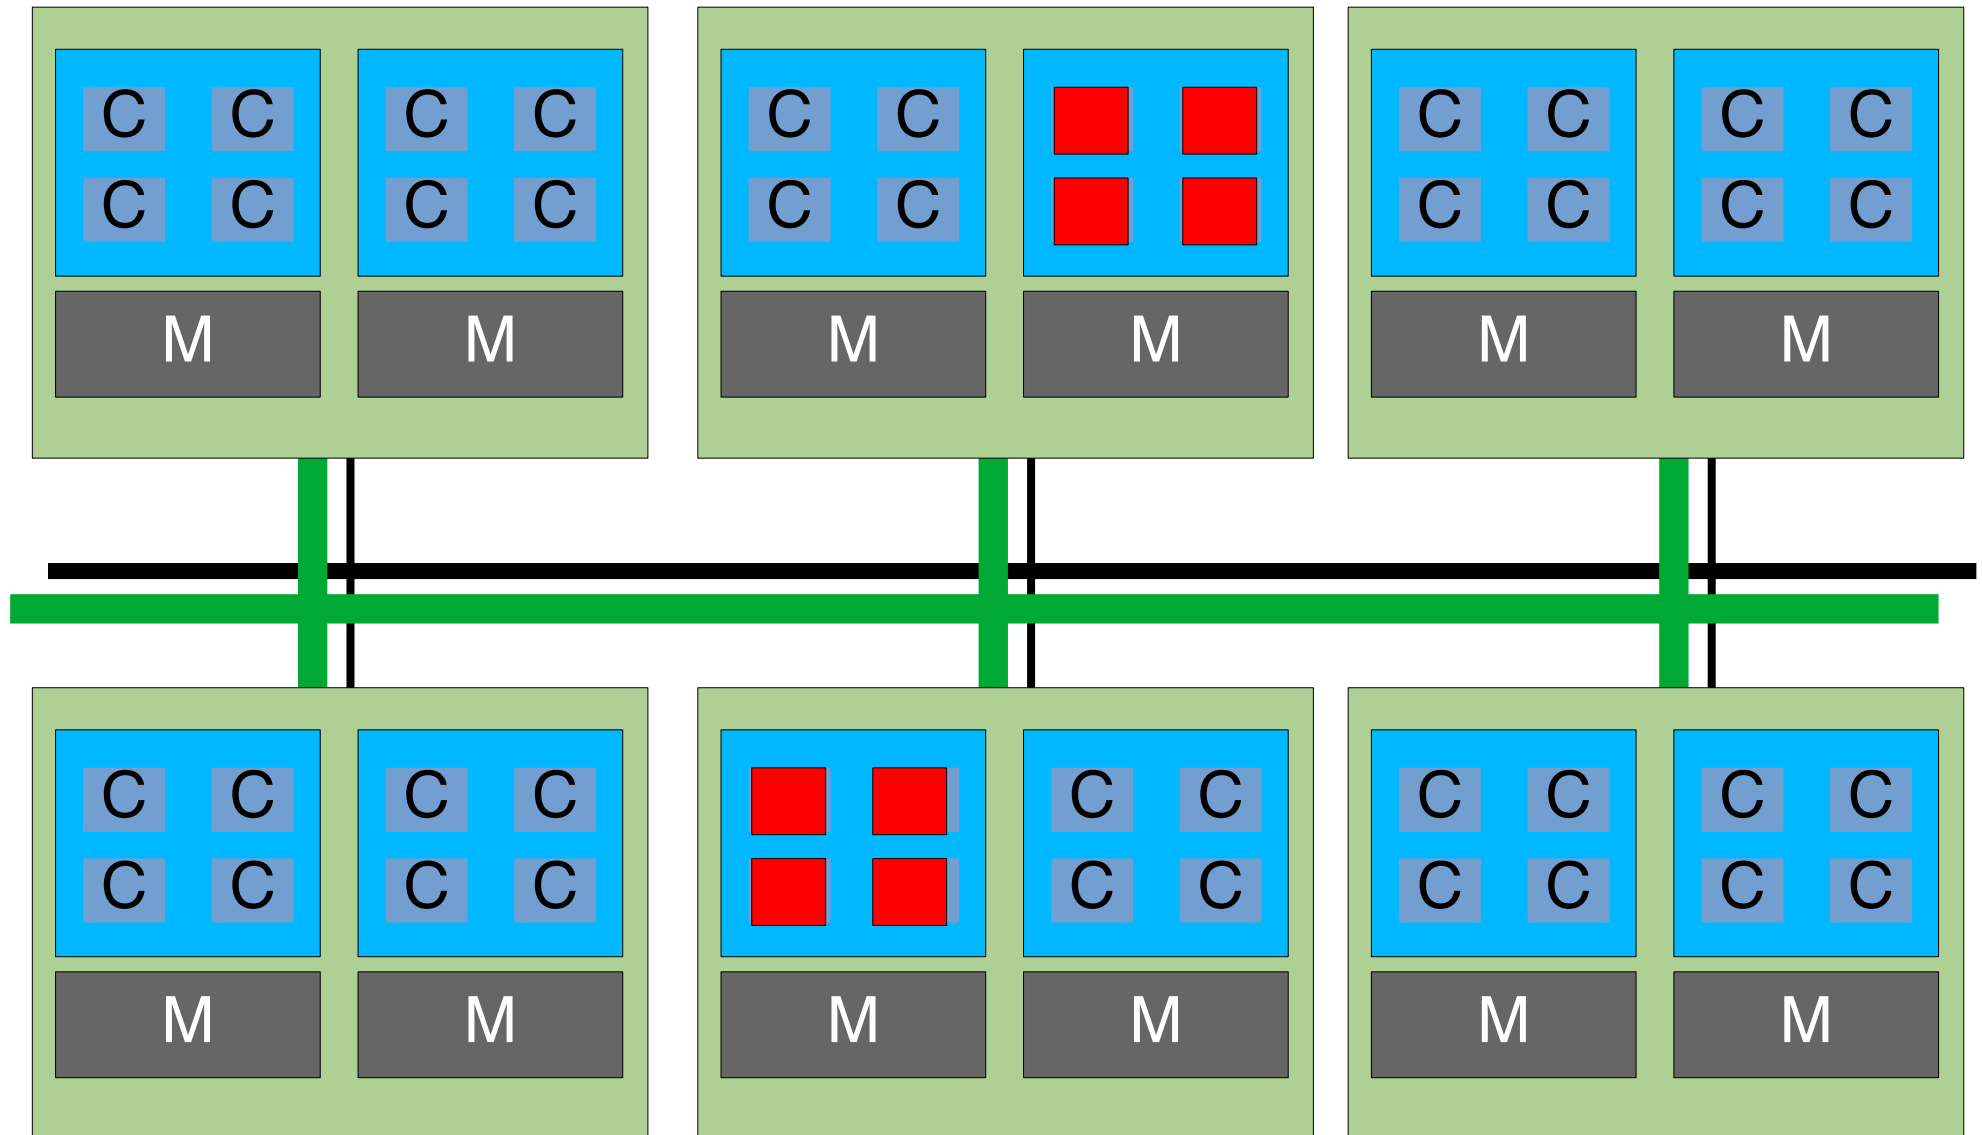
```
#BSUB -J MPIapp  
#BSUB -o MPIapp_%J.out  
#BSUB -n 8  
#BSUB -R "span[ptile=4]"
```

```
module load mpi/<version>
```

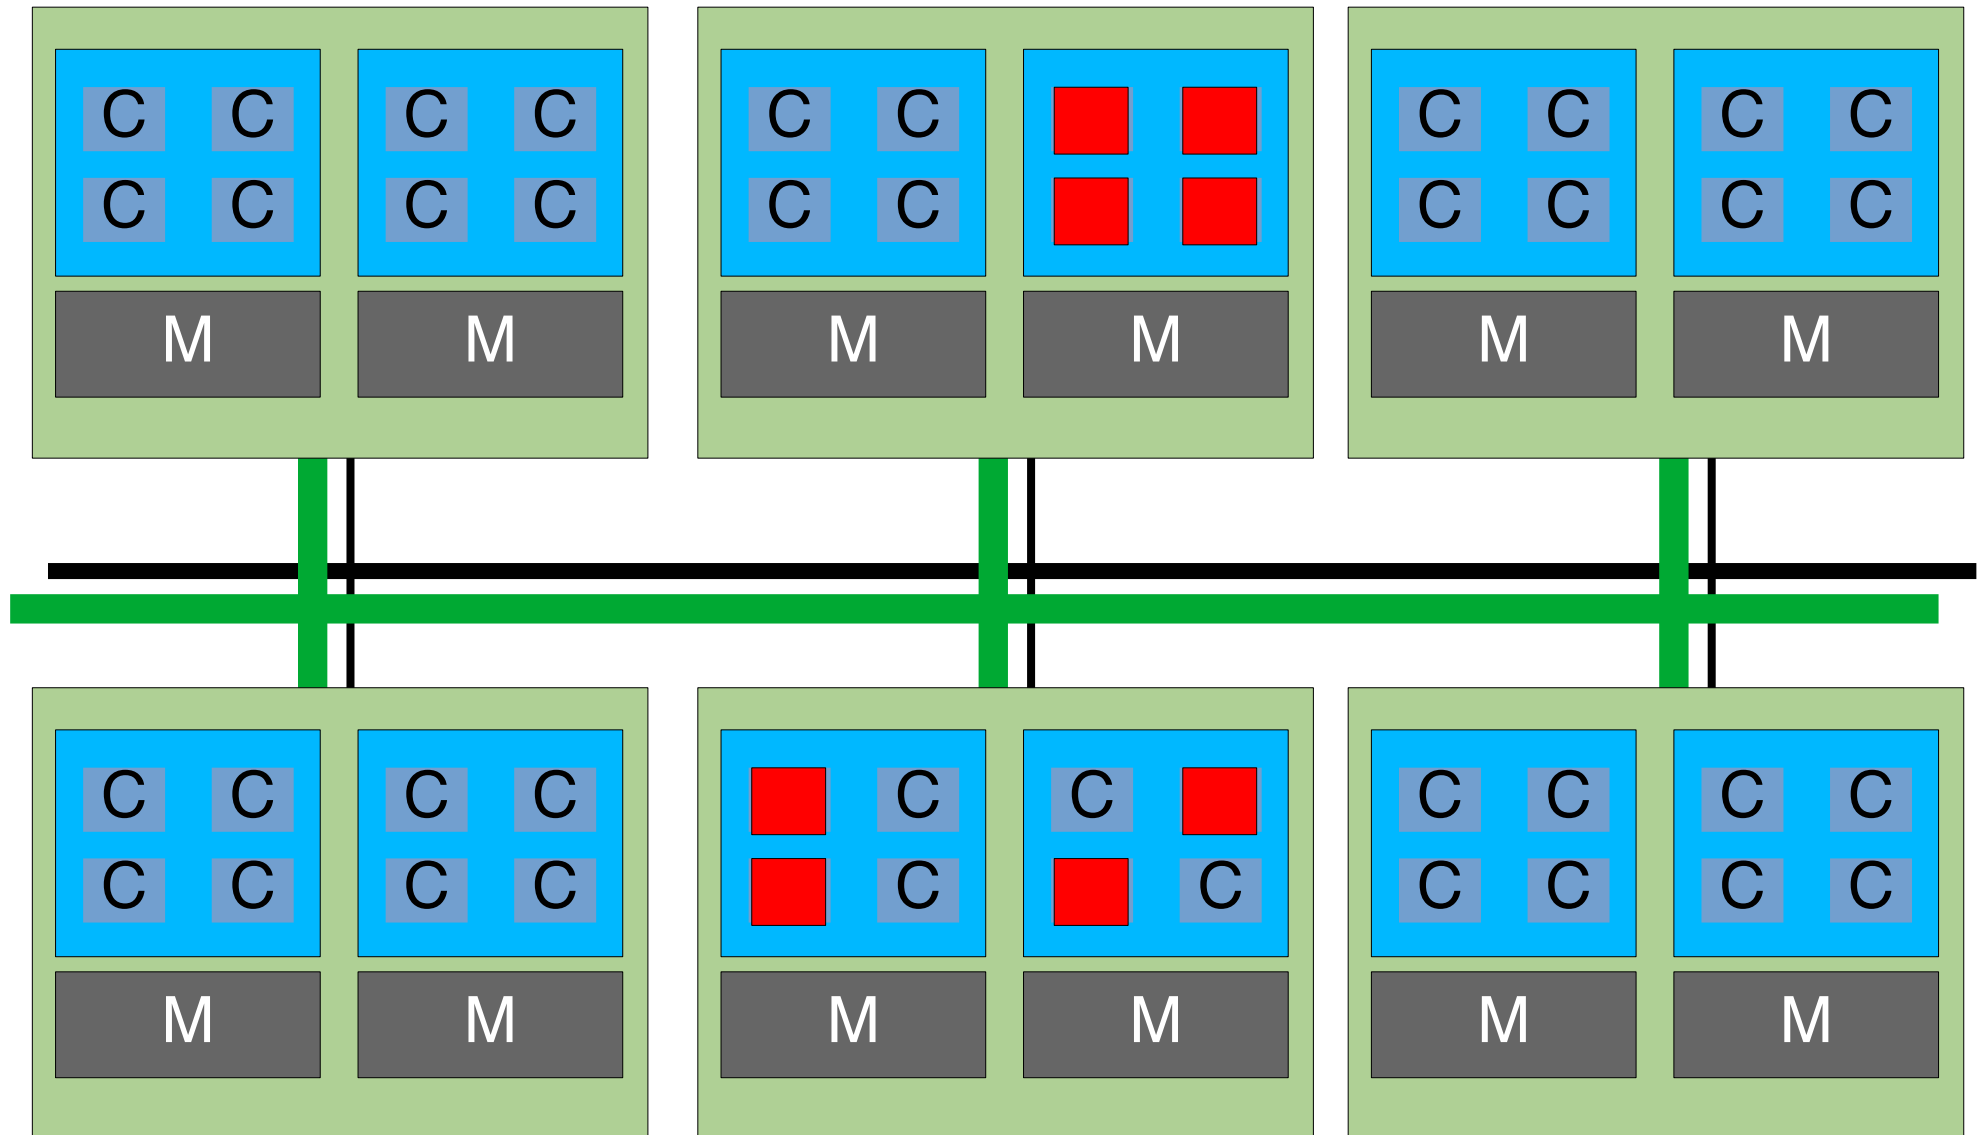
```
mpirun ./myapp
```

- ❑ What do we get?

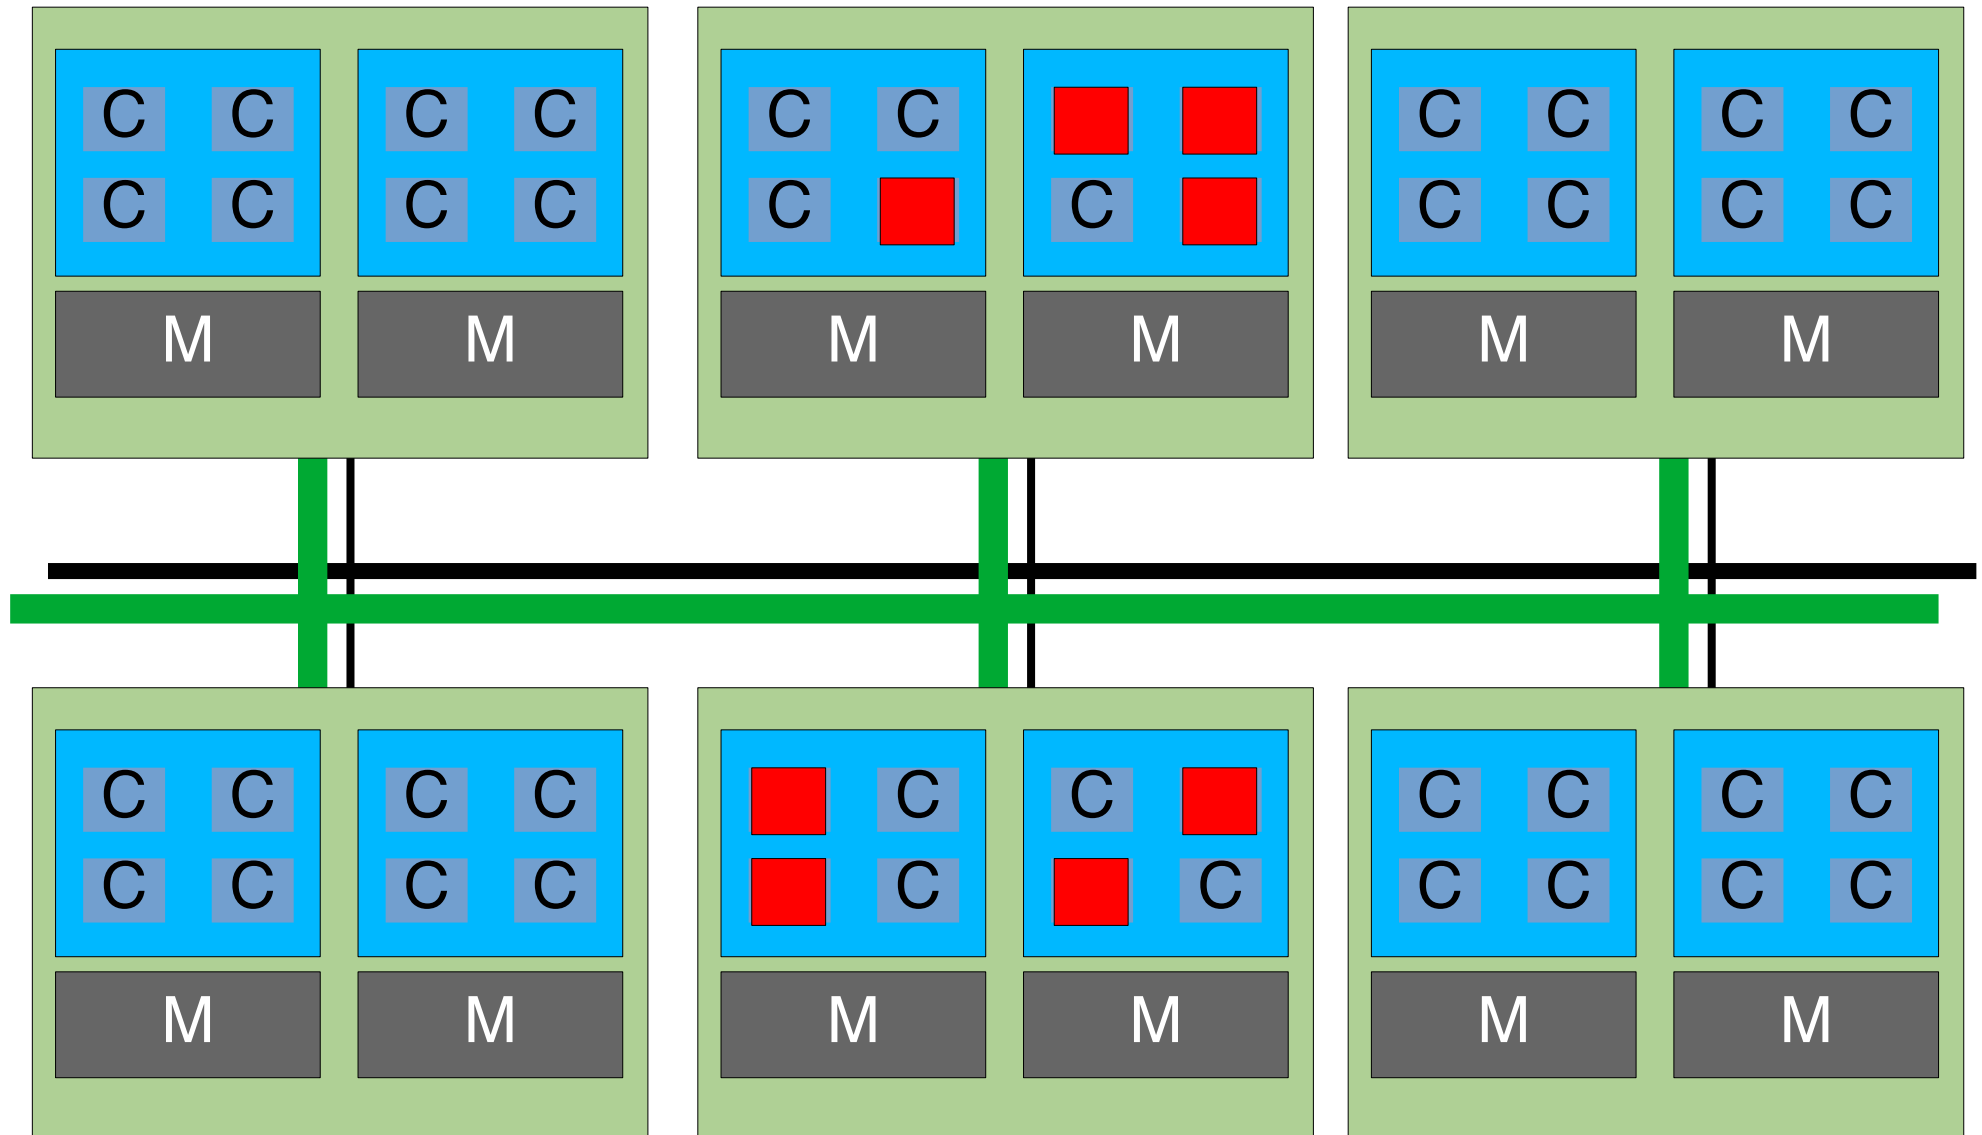
The “real” world



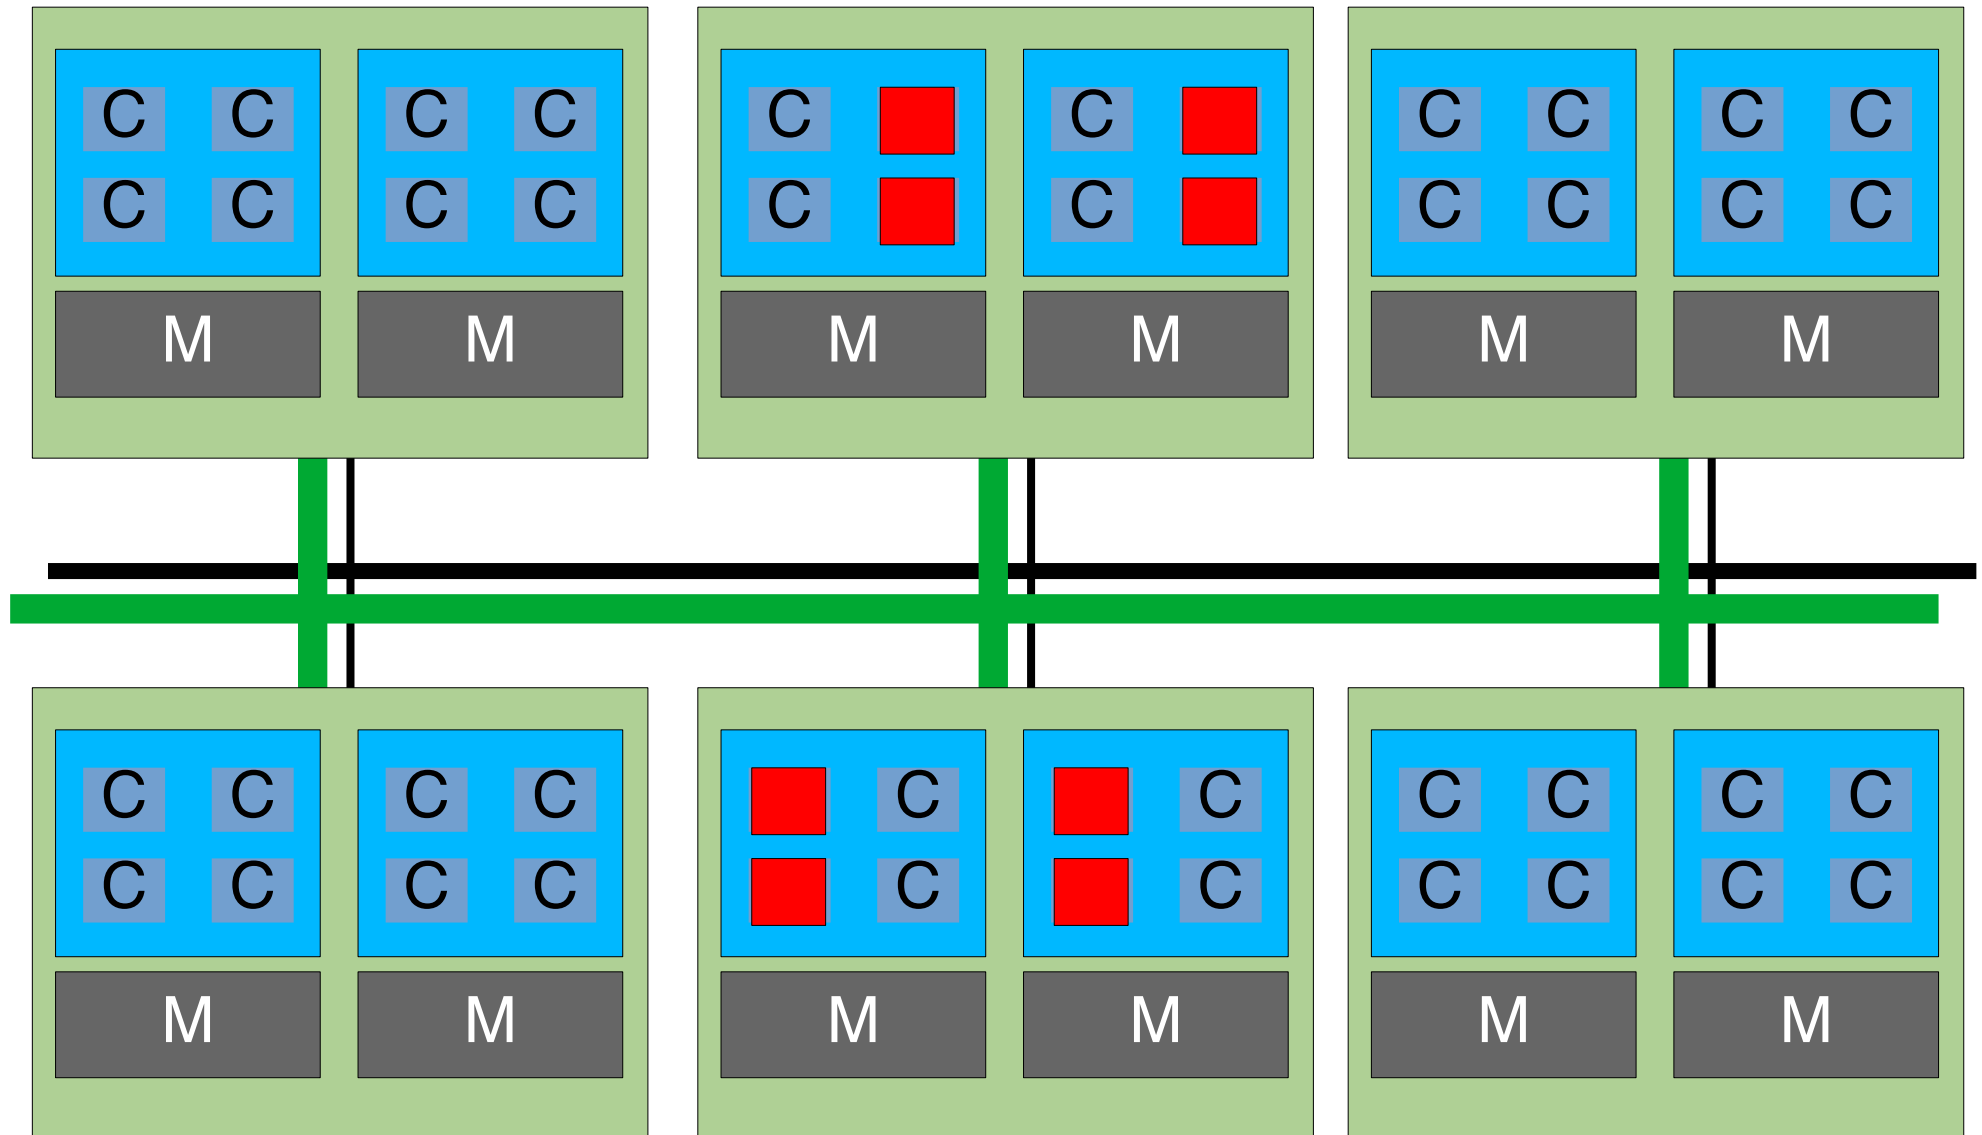
The “real” world



The “real” world



The “real” world



The “real” world

Scenario 3:

- ❑ use 4 nodes, 2 ranks per node

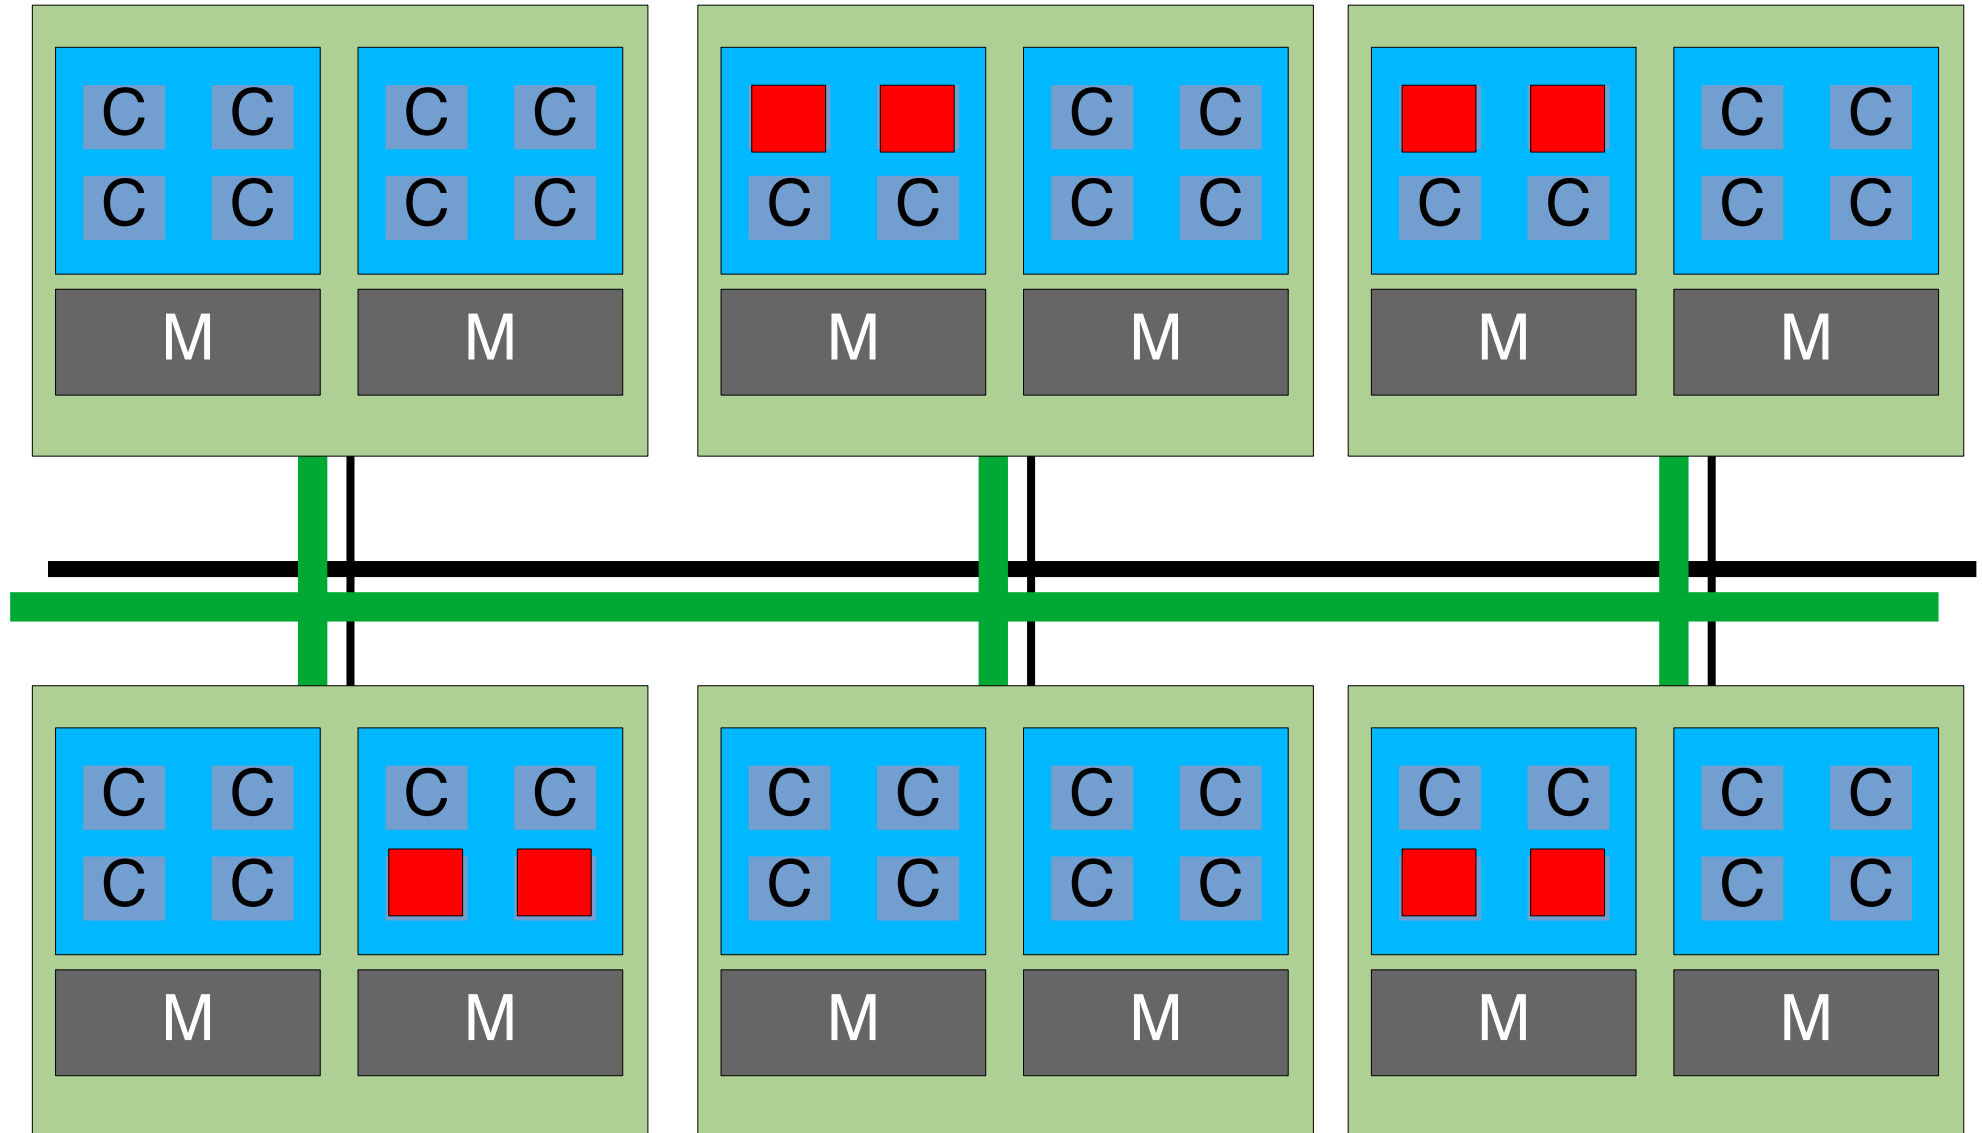
```
#BSUB -J MPIapp  
#BSUB -o MPIapp_%J.out  
#BSUB -n 8  
#BSUB -R "span[ptile=2]"
```

```
module load mpi/<version>
```

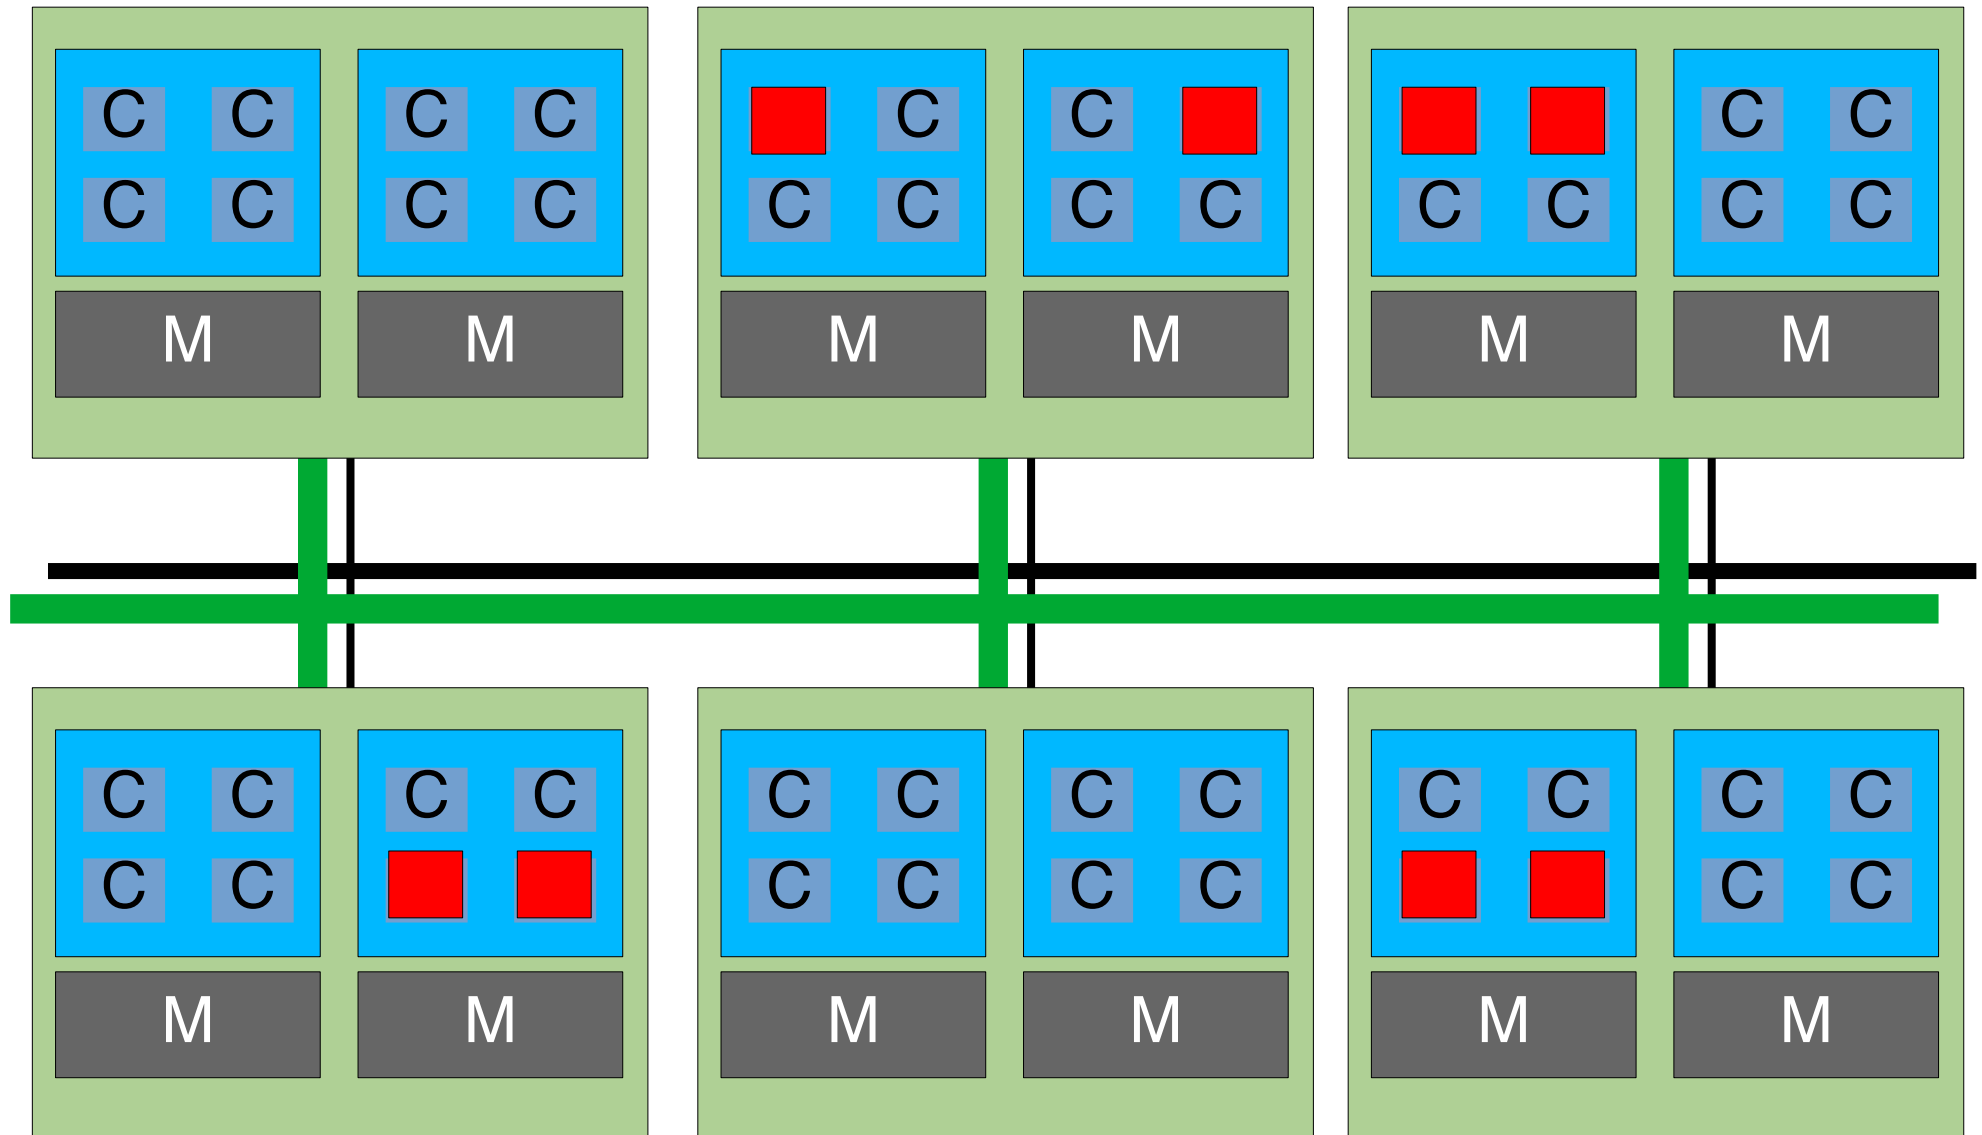
```
mpirun ./myapp
```

- ❑ What do we get?

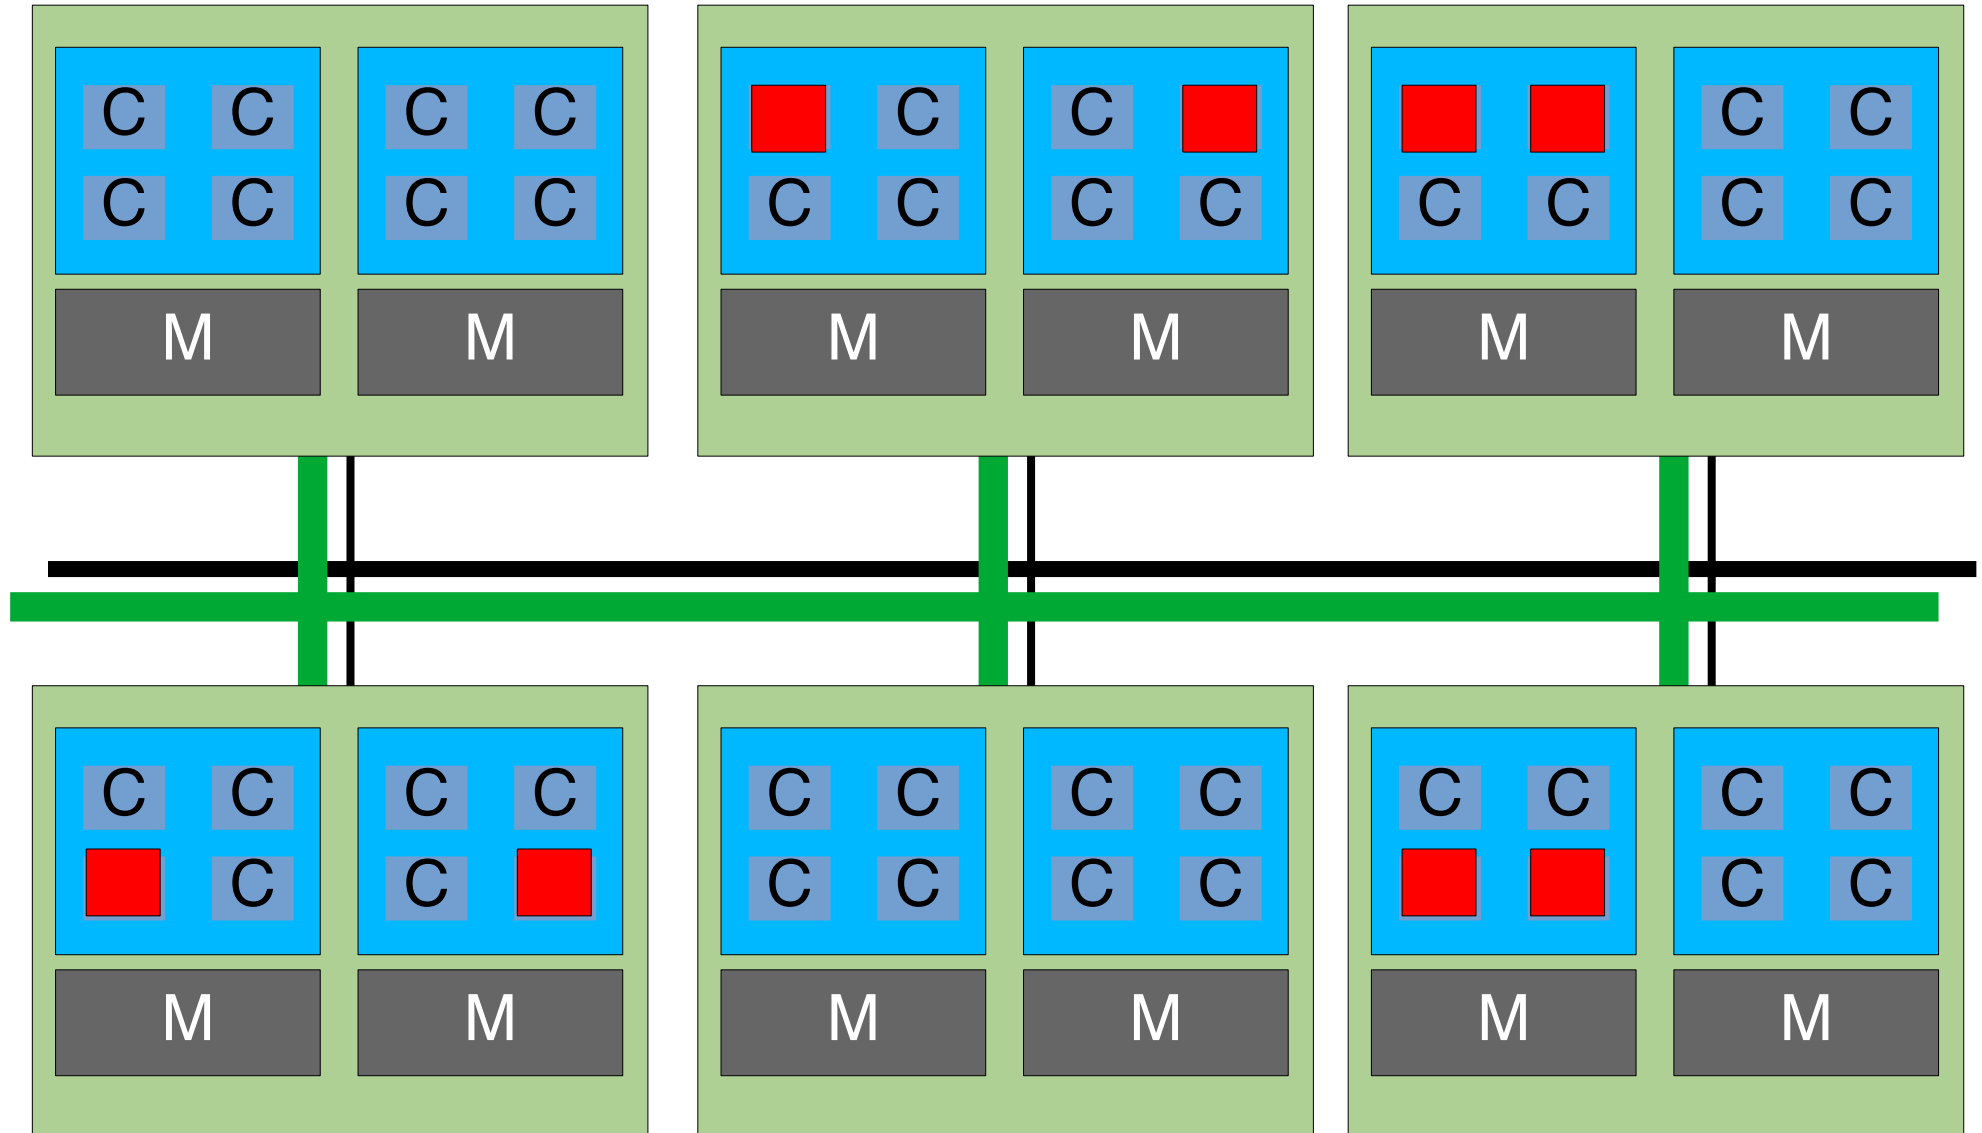
The “real” world



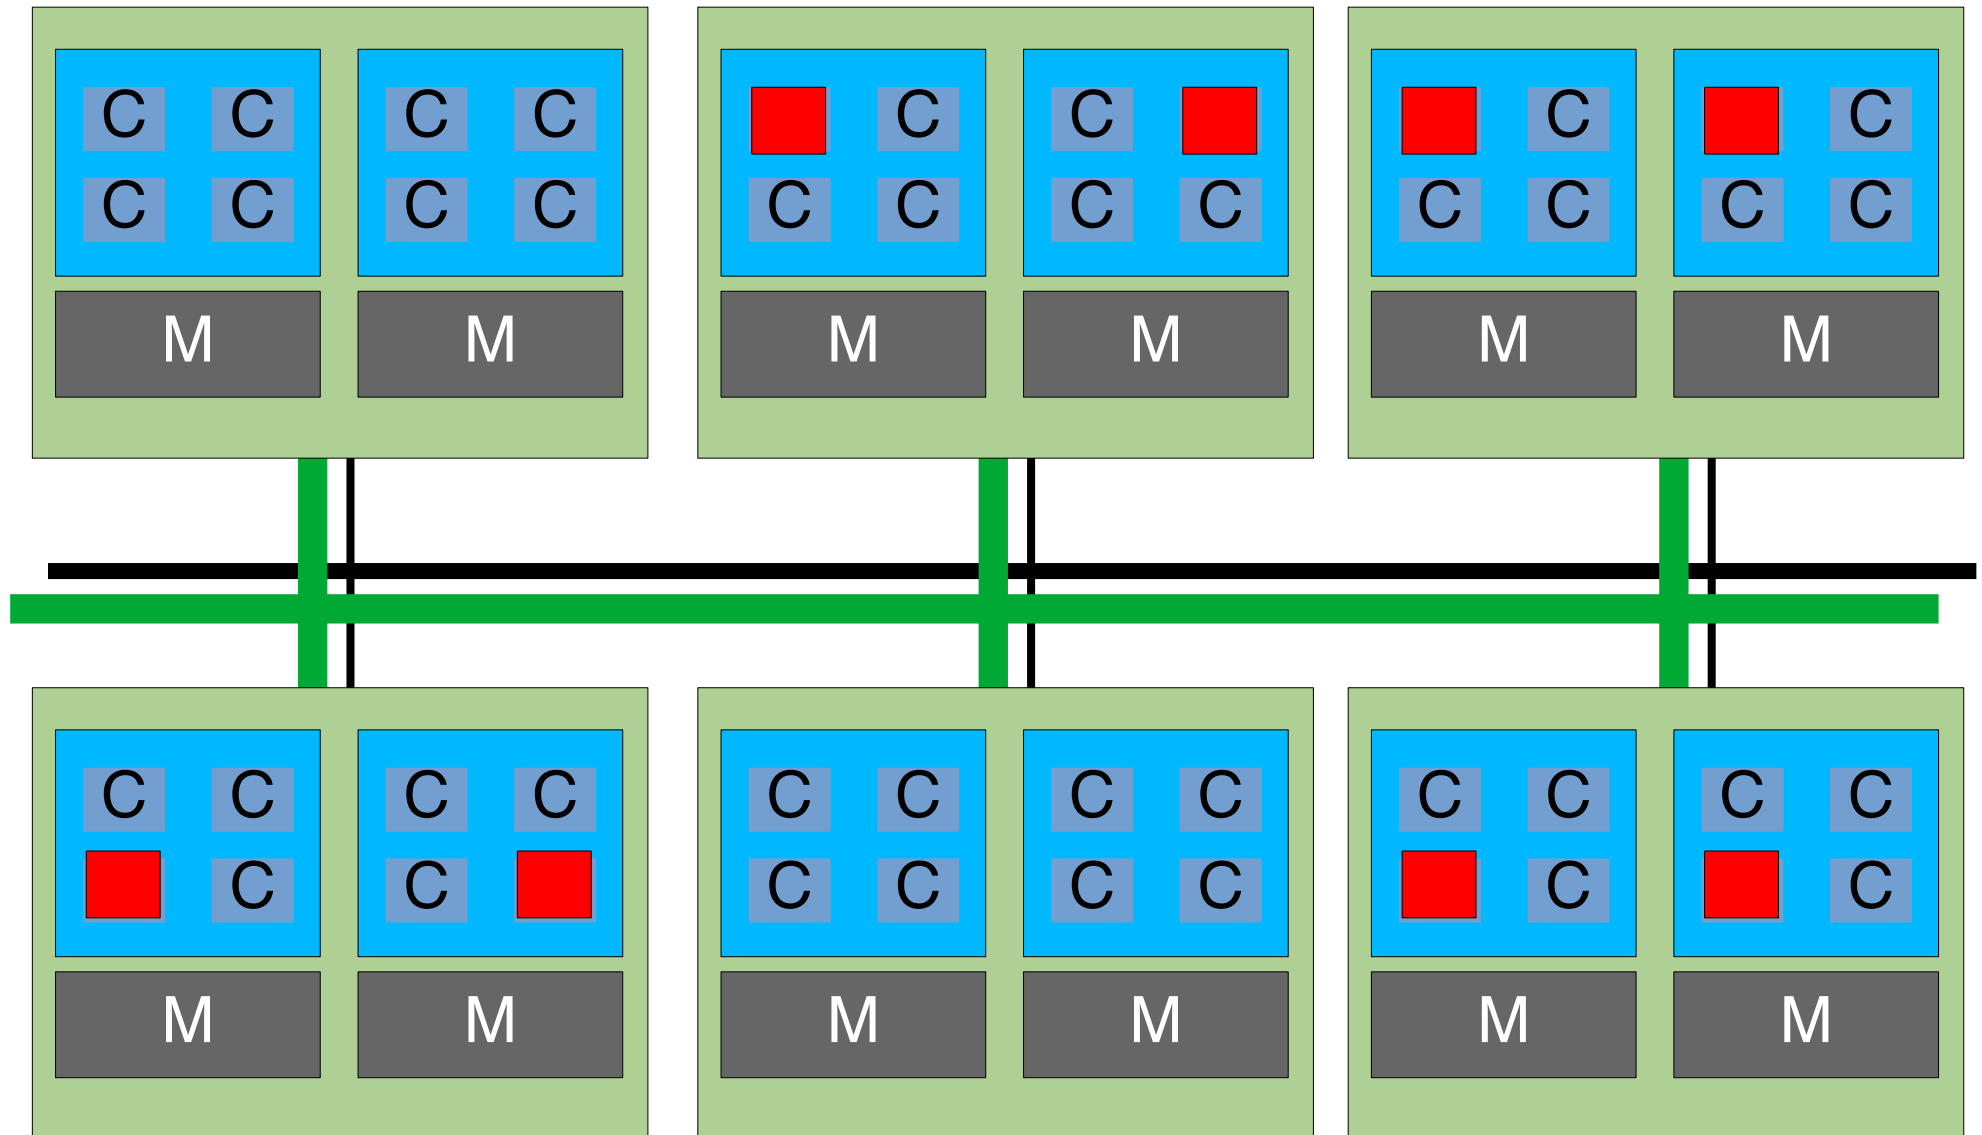
The “real” world



The “real” world



The “real” world



The “real” world

- ❑ If our MPI application needs a lot of memory bandwidth, then we want to distribute the ranks over as many CPU sockets as possible.
- ❑ Here, the OpenMPI runtime options for mapping can help
- ❑ Syntax:
 - ❑ `mpirun -map-by ppr:n:<resource>`
 - ❑ ppr stands for “process(es) per resource”
 - ❑ `<resource>` can be package (a CPU socket), core, ...
 - ❑ check the documentation for more details

The “real” world

Scenario 3 – with mapping

- ❑ use 4 (full) nodes, 2 ranks per node

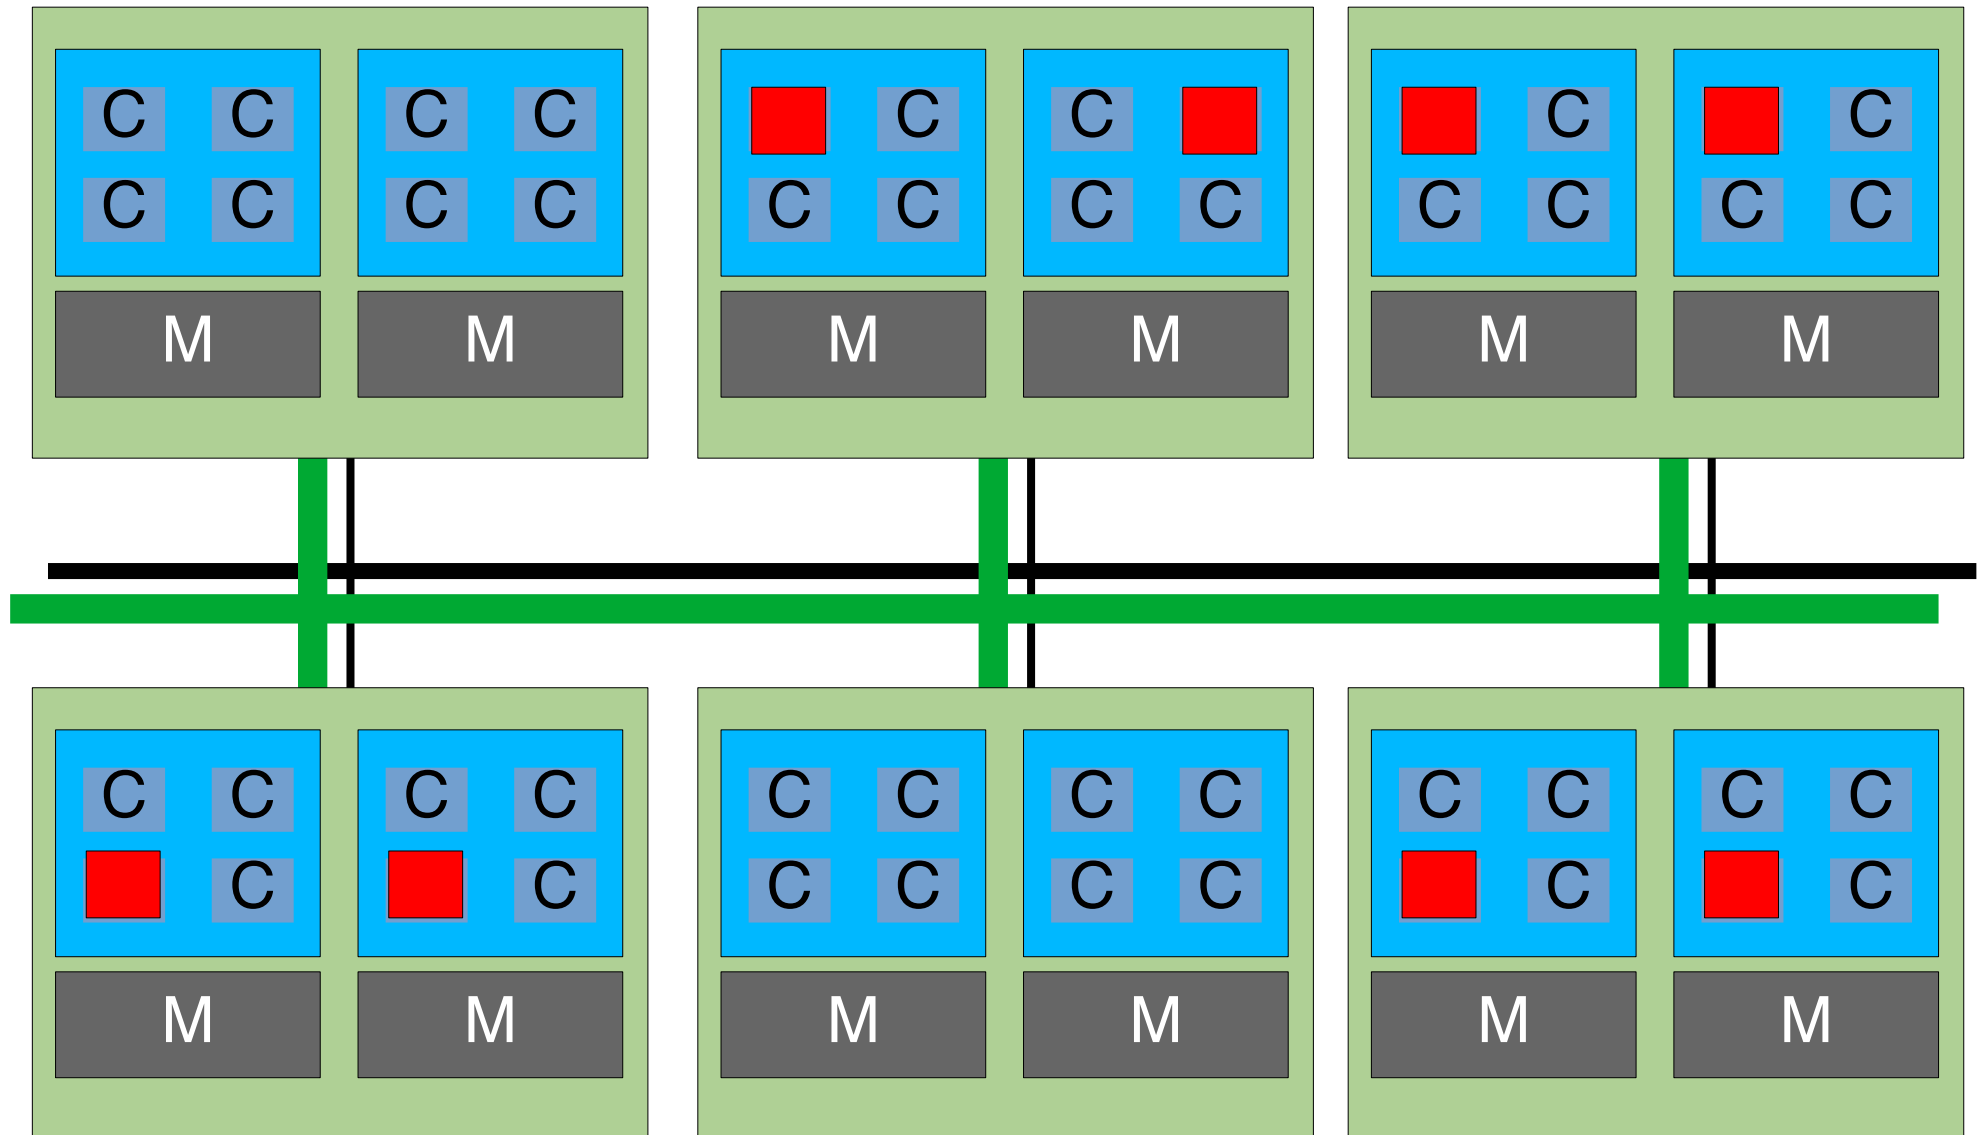
```
#BSUB -J MPIapp
#BSUB -o MPIapp_%J.out
#BSUB -n 32
#BSUB -R "span[ptile=8]"
```

```
module load mpi/<version>
```

```
mpirun -np 8 -map-by ppr:1:package ./myapp
```

- ❑ Then we get something like ... on the next slide

The “real” world



The “real” world

- ❑ The solution above gives us the maximum bandwidth for our problem
- ❑ ... for the price of allocating more resources than we need!
- ❑ There are other ways to implement this, e.g. using the affinity options of LSF, both with and without additional MPI runtime options
 - ❑ ... helps to save resources
 - ❑ ... but that’s beyond the scope of today!
- ❑ Today’s lab: measure runtime of such a placement scenario