

02616

Large-scale Modelling

a look into Python

Python multiprocessing

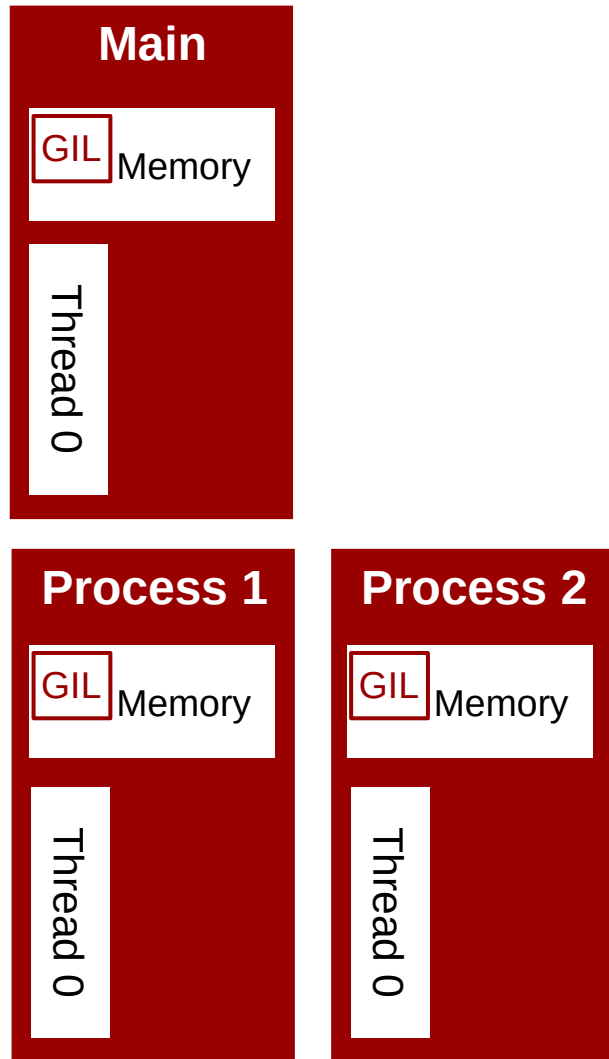
- Bypasses GIL problems (synchronization locks)
 - Each process has its own memory space, the GIL ends
 - multiprocessing is low-level
 - concurrent.futures.* is the high-level
 - We'll only do low-level!
- All code runs in parallel
- Good for CPU-bound tasks

```
with multiprocessing.Pool(4) as pool:
    rets = pool.map(hard_work, range(20))
```

Calls `hard_work` 20 times, with arguments 0..19
(think `np.array_split`!)

- Can have substantial overhead!

How multiprocessing works



Pro:

Processes can always run in parallel

Cons:

Processes have more overhead

No shared memory – must explicitly copy

Launching processes

Pool:

```
with mp.Pool(4) as pool:
```

Pros:

- Simple to use
- Context cleans up after use
- Easily allowing return

Cons:

- Hard to customize for un-even tasks
- With tasks > NCPU's it can be hard to use communication channels

Process:

```
p1 = mp.Process(...)  
p2 = mp.Process(...)
```

Pros:

- Full control
- Simple communication through manual configuration

Cons:

- Requires manual work for return values
- Manual context clean-up

Python communication

- Communication is *hard*
- Using Queue
Simple, but non-deterministic.

```
with mp.Pool(4) as pool:  
    queue = mp.Manager().Queue()  
    pool.map(run, [queue for _ in range(20)])  
while not queue.empty():  
    print(queue.get())
```

- Using Pipe
Explicit communication between two processes
(similar to MPI_Send/MPI_Recv)
Always *blocking*!

We'll see examples in today's exercises!

Overhead

- What does it cost to launch additional processes?

```
def hard_work(a):
    ...
```

```
with mp.Pool(1) as pool:
    pool.map(hard_work, [1])
```

```
queue = mp.Queue()
p = mp.Process(target=wrap,
               args=(queue,))
p.start()
queue.get()
```

```
% time python overhead.py
```

Pool

Process

Walltime
(total)

No multiprocessing

```
real 1.27
user 1.19
sys 0.14
```

CPU time in system
calls

```
real 1.10
user 1.08
sys 0.01
```

CPU time in user
process

```
real 1.20
user 1.10
sys 0.11
```

```
real 0.63
user 3.29
sys 0.05
CPU 530%
```

Bandwidth

Transport of messages

1. Start timing
2. Send a message
3. Stop timing once you know it has been received

We'll come back to why this is not always easy!

Bandwidth = [message size] / [time spent]

Carefully think about how you time!

- Where should you start timing?
- Where should you end timing?
- Once you have full time, what does it *actually* time?
- How should you execute your program?

Everything uses resources!

Bandwidth

```
t0 = time()  
<initialize program>  
t1 = time()  
<do heavy computations>  
t2 = time()  
<finalize program>  
t3 = time()
```

6(7) different timings!

Bandwidth & latency!

Measuring bandwidth and latency is difficult!

Bandwidth

Can mean:

- 1) time of communication, excluding startup delay
- 2) time of communication, including startup delay

Be specific when using this term!

Note units! (Mb vs. MB)

Latency

Can mean:

- 1) the startup delay, excluding transfer
- 2) the communication delay, including transfer

Be specific when using this term!

$$T(s) = \mu_s + \frac{1}{B}s$$

Today's exercises

- Play with multiprocessing
- Gathering of results in a single process
- Reduction algorithms
- A bandwidth plot (be careful!)

The simplest is to use `Pool`, so start with that.
Then implement with `Process` (if you want!)

Volunteers for next week?