

02616

# Large-scale Modelling

one library to rule MPI

`mpi4py`

# MPI

Welcome to MPI the fundamental Message Passing Interface.

MPI is a *standard* with official C and fortran bindings

**All** other language API's are **NOT** official!

including mpi4py!

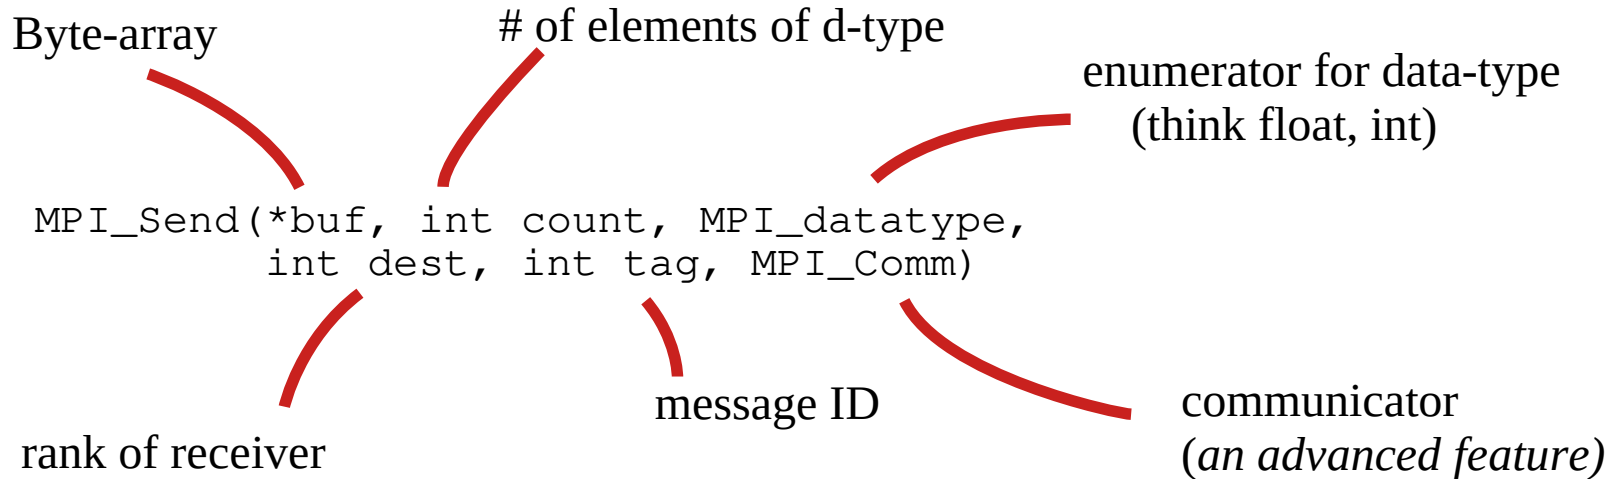
MPI implementations:

- OpenMPI
- MPICH
- IntelMPI (MPICH)
- ... (vendors)

used extensively!

<https://mpi4py.readthedocs.io/en/stable/reference/mpi4py.MPI.Comm.html>  
<https://docs.open-mpi.org/en/v5.0.x/man-openmpi/man3/index.html>

# MPI – behind the scenes



`MPI_Rsend` ready-send (Recv *must* have been issued!)  
`MPI_Bsend` buffered send (manual setup of a buffer in MPI)  
`MPI_Ssend` synchronous send (blocking)  
`MPI_Send` implementation specific *standard* send (typically Bsend or Ssend)  
`MPI_I?send` non-blocking variants!

Any  
combination

```

MPI_Recv(*buf, int max_count, MPI_datatype,
        int source, int tag, MPI_Comm, *MPI_Status)
MPI_Irecv(*buf, int max_count, MPI_datatype,
        int source, int tag, MPI_Comm, *MPI_Request)
  
```

# MPI → mpi4py

mpi4py documentation is, ..., lacking...  
(an open-source project ready for huge impact from contributors!)

`MPI_X(..., MPI_Comm, ...)` → `mpi4py.Comm.X(...)`

Typically, there are *sane* defaults:

- `count` defaults to number of elements in array
- `datatype` defaults to array type
- `source` defaults to `MPI_ANY_SOURCE`
- `tag` defaults to `MPI_ANY_TAG | 0`

When in doubt, look at the documentation of both  
OpenMPI + mpi4py!

# mpi4py – the basics

**Send**(buf, dest, tag=0)

Blocking send.

```
int MPI_Send(const void *buf, int count, MPI_Datatype datatype, int dest,
             int tag, MPI_Comm comm)
```

## Note

This function may block until the message is received. Whether **Send** blocks or not depends on several factors and is implementation dependent.

- Parameters:
- buf (**BufSpec**)
  - dest (*int*)
  - tag (*int*)

Return type: *None*

**Recv**(buf, source=ANY\_SOURCE, tag=ANY\_TAG, status=None)

Blocking receive.

```
int MPI_Recv(void *buf, int count, MPI_Datatype datatype,
             int source, int tag, MPI_Comm comm, MPI_Status *status)
```

## Note

This function blocks until the message is received.

- Parameters:
- buf (**BufSpec**)
  - source (*int*)
  - tag (*int*)
  - status (*Status* | *None*)

Return type: *None*

# mpi4py – the basics

```
from mpi4py import MPI
C = MPI.COMM_WORLD
a = np.array([1., 2.])
```

## mpi4py.typing.BufSpec

Buffer specification.

- `Buffer`
- `Tuple[ Buffer, Count ]`
- `Tuple[ Buffer, TypeSpec ]`
- `Tuple[ Buffer, Count, TypeSpec ]`
- `Tuple[ Bottom, Count, Datatype ]`

- `C.Send(a, 1)`
- `C.Send((a, 1), 1)`
- `C.Send((a, MPI.DOUBLE_PRECISION), 1)`
- `C.Send((a, 2, MPI.DOUBLE_PRECISION), 1)`

# mpi4py – the basics

`Send(buf, dest, tag=0)`

Blocking send.

## Note

This function may block until the message is received. Whether `Send` blocks or not depends on several factors and is implementation dependent.

Parameters:

- `buf` (`BufSpec`)
- `dest` (`int`)
- `tag` (`int`)

Return type: `None`

`send(obj, dest, tag=0)`

Send in standard mode.

Parameters:

- `obj` (`Any`)
- `dest` (`int`)
- `tag` (`int`)

Return type: `None`

`recv(buf=None, source=ANY_SOURCE, tag=ANY_TAG, status=None)`

Receive.

Parameters:

- `buf` (`Buffer` | `None`)
- `source` (`int`)
- `tag` (`int`)
- `status` (`Status` | `None`)

Return type: `Any`

Has implications!

Use only for non-critical sections!

# MPI – behind the scenes

```
MPI_Recv(*buf, int max_count, MPI_datatype,  
         int source, int tag, MPI_Comm, *MPI_Status)
```

MPI\_Recv is *magic*:

- Buffer size  $\geq$  sent message
- Data-type needs not be the same (MPI\_Send: float + MPI\_Recv: complex)
- Can receive from arbitrary source (MPI\_ANY\_SOURCE)
- Can receive any message ID (MPI\_ANY\_TAG)
- The MPI\_Status is your gateway to information!

Only use a status if you:

- Use MPI\_ANY\_\*, or
- Are not sure of how many bytes received
- *Check API documentations for details of data-extraction!*
  - MPI\_Get\_count / status.Get\_count
  - status.source or status.tag

Else use MPI\_STATUS\_IGNORE (the default)!



# MPI non-blocking

- Allows overlapping *communication* and *computation*!
- Requires manual completion!

```
Comm.Isend(..., request=request)
# do computation
Comm.wait(request)
```

- All non-blocking variants returns a *request* handle (not a status!)
- Completes only after one of the following succeeds:

- |               |   |                              |
|---------------|---|------------------------------|
| • MPI_Test    | } | <i>May return true/false</i> |
| • MPI_Testany |   |                              |
| • MPI_Testall |   |                              |
| • MPI_Wait    | } | <i>Ensures completion!</i>   |
| • MPI_Waitany |   |                              |
| • MPI_Waitall |   |                              |

They all return a `MPI_Status` (or list thereof).

see sample.py

# MPI collectives

Collectives are messages done with *all* the ranks participating in the distribution.

Collectives:

- `Bcast` a value to all ranks (one  $\rightarrow$  all)
- `Scatter` values from a list to all ranks (one  $\rightarrow$  all)
- `Gather` values into a list on a single rank (all  $\rightarrow$  one)
- `Reduce` values on a single rank (all  $\rightarrow$  one)

Global collectives:

- `Allgather` equivalent to `Gather` + `Bcast`
- `Allreduce` equivalent to `Reduce` + `Bcast`
- `Alltoall?` equivalent to all calling `Scatter/Gather` for every rank

# MPI sample code

Let's go through (some of) it together!

Volunteers for next week?